SchemaOnRead: A Package for Schema-on-Read in R

by Michael J. North

Abstract **SchemaOnRead** is a CRAN package that provides an extensible mechanism for importing a wide range of file types into R as well as support for the emerging schema-on-read paradigm in R. The schema-on-read tools within the package include a single function call that recursively reads folders with text, comma separated value, raster image, R data, HDF5, NetCDF, spreadsheet, Weka, Epi Info, Pajek network, R network, HTML, SPSS, Systat, and Stata files. It also recursively reads folders (e.g., schemaOnRead("folder")), returning a nested list of the contained elements. The provided tools can be used as-is or easily customized to implement tool chains in R. This paper's contribution is that it introduces and describes the **SchemaOnRead** package and compares it to related R packages.

Introduction

SchemaOnRead is a CRAN package that provides an extensible mechanism for importing a wide range of file types into R as well as support for the emerging schema-on-read paradigm in R. The tools within the package include a single function call (e.g., schemaOnRead("filename")) that reads text (TXT), comma separated value (CSV), raster image (BMP, PNG, GIF, TIFF, and JPG)[1], R data (RDS), HDF5, NetCDF, spreadsheet (XLS, XLSX, ODS, and DIF), Weka Attribute-Relation File Format (ARFF), Epi Info (EPIINFO), Pajek network (NET), R network (PAJ), HTML, SPSS (SAV), Systat (SYS), and Stata (DTA) files. It also recursively reads folders (e.g., schemaOnRead("folder")), returning a nested list of the contained elements. The provided tools can be used as-is or easily customized to implement tool chains in R. This paper's contribution is that it introduces and describes the **SchemaOnRead** package and compares it to related R packages. In the sections that follow, this paper presents usage examples, discusses user defined processors, reviews the related work, explains the origin of the package name, summarizes the package contents, and then provides concluding thoughts.

Examples

A simple way to use SchemaOnRead is to conveniently load a file without needing to handle the specifics of the file format. In this case the result is a variable containing the file contents. Individual files can also be easily accessed without needing to known the specifics of the file format as shown in Figure 1. The file contents can be accessed using the xmlFile variable. All of the source code and example data can be found at https://github.com/drmichaelnorth/SchemaOnRead.

```
library(SchemaOnRead)
xmlFile <- schemaOnRead("../inst/extdata/data.xml")
```

Figure 1: Loading a File with SchemaOnRead

Another way to use SchemaOnRead is to recursively load a folder. The result is a named list of elements for each entry in the folder's tree as shown in Figure 2. Sub-elements (e.g., files or subfolders) of a folder can be accessed using the R named list ($) operator followed by the sub-element name. An example showing how to read a folder tree starting in '../inst/extdata' is shown in Figure 3. In this case, the contents of the 'dir1/Data.csv' file within '../inst/extdata' is shown by accessing 'results$dir1$Data.csv' as needed. The path also provides the data provenance. Files or folders with names that do not conform to standard R variable naming requirements can be accessed using single quote notation (e.g., results$'Nonconforming Name'). The resulting named list notation also provides convenient access to files and folders using integrated development environments for R that support automatic code completion. An RStudio (RStudio, 2015) example is shown in Figure 4.

---

[1]Image processing applications are becoming increasingly popular for purposes such as pattern recognition and machine vision. These applications often read large numbers of files during their training and testing phases. Image file import has been added to **SchemaOnRead** to support this use case.
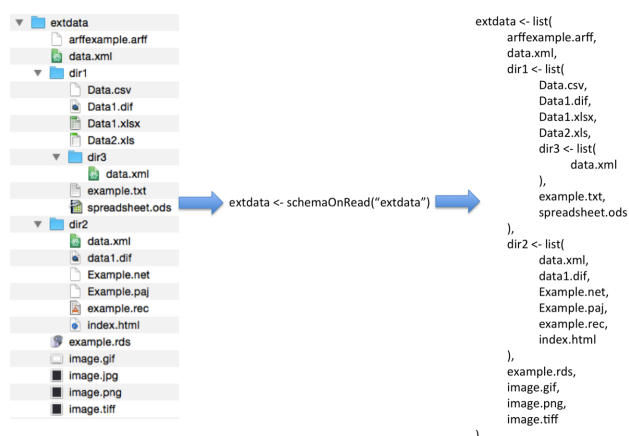
Figure 2: Reading a Nested Set of Folders

```
library(SchemaOnRead)
results <- schemaOnRead("../inst/extdata")
print(results$dir1$Data.csv)
```

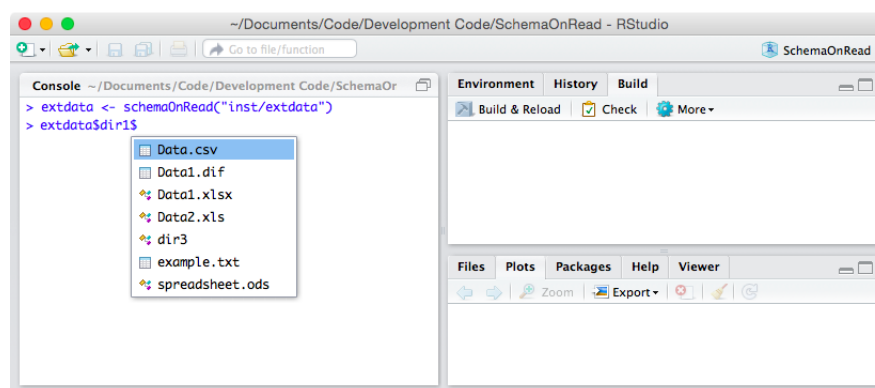**Figure 3**: How to Access Data from a File Nested within Folders

Figure 4: Using SchemaOnRead for Convenient Access to Files and Folders in RStudio

The SchemaOnRead verbose flag can be used to trace a call's progress or diagnose issues as shown in Figure 5. Example results on the R console are shown in Figure 6.

```
library(SchemaOnRead)
folder <- schemaOnRead("../inst/extdata", verbose = TRUE)
```

Figure 5: Use of the SchemaOnRead verbose Flag

```
schemaOnRead processing ../inst/extdata
schemaOnRead processing ../inst/extdata/arffexample.arff
schemaOnRead processing ../inst/extdata/data.xml
schemaOnRead processing ../inst/extdata/dir1
schemaOnRead processing ../inst/extdata/dir1/Data.csv
schemaOnRead processing ../inst/extdata/dir1/Data1.dif
schemaOnRead processing ../inst/extdata/dir1/Data1.xlsx
schemaOnRead processing ../inst/extdata/dir1/Data2.xls
schemaOnRead processing ../inst/extdata/dir1/dir3
schemaOnRead processing ../inst/extdata/dir1/dir3/data.xml
schemaOnRead processing ../inst/extdata/dir1/example.txt
schemaOnRead processing ../inst/extdata/dir1/spreadsheet.ods
schemaOnRead processing ../inst/extdata/dir2
schemaOnRead processing ../inst/extdata/dir2/data.xml
```

Figure 6: Console Results of the SchemaOnRead verbose Flag

User Defined Processors

New processors can be defined to support user-specified processing. New processors are normally prepended to the front of the default list to allow them to take precedence while still allowing the standard processors to work if needed. Alternatively, a list of processors that just recursively scans folders can be found by calling the schemaOnReadSimpleProcessors function. User-specified processors can be added to this list to create a fully customized tool chain. An example showing how to create a simple files processor is given in Figure 7. A more detailed example of a Microsoft Excel spreadsheet processor is shown in Figure 8.

```
## Load the needed library.
library(SchemaOnRead)

## Define a new processor.
newProcessor <- function(path, ...) {

  # Check the file existance and extensions.
  if (!SchemaOnRead::checkExtensions(path, c("xyz"))) return(NULL)

  ## As an example, attempt to read an XYZ file as a CSV file.
  read.csv(path, header = FALSE)

}

## Define a new processors list.
newProcessors <- c(newProcessor, SchemaOnRead::defaultProcessors())

# Use the new processors list.
schemaOnRead(path = "../inst/extdata", processors = newProcessors)
```

Figure 7: Creating and Using a Simple User-Defined Processor

```
## Load the needed library.
library(SchemaOnRead)

## Define a new processor.
newSpreadsheetProcessor <- function(filePath = ".", ...) {

  # Check the file existance and extensions.
  if (!SchemaOnRead::checkExtensions(filePath, c("xls", "xlsx"))) return(NULL)

  # Read the workbook's worksheet names.
  worksheets <- readxl::excel_sheets(filePath)

  # Read the workbook's worksheets.
  workbook <- lapply(worksheets, readxl::read_excel, path = filePath)

  # Name the worksheets.
  names(workbook) <- worksheets

  # Return the results.
  workbook

}

## Define a new processors list.
newProcessors <- c(newSpreadsheetProcessor, SchemaOnRead::defaultProcessors())

# Use the new processors list.
schemaOnRead(path = "../inst/extdata", processors = newProcessors)
```

Figure 8: Creating and Using a User-Defined Processor for Spreadsheets

Related Work

Several R packages provide support for importing diverse file formats into R. Examples include rio, readbitmap, and foreign.

The **rio** package (Chan et al., 2015) is the closest in functionality to **SchemaOnRead**. **rio** provides file reading functions for a wide range of formats including text files, fixed format files, spreadsheet files (XLS, XLSX, ODS, and DIF),Stata, JSON, SPSS, Weka, Epi Info, serialized R objects, saved R objects, SAS, Minitab, Systat, shallow XML files, FORTRAN data files, and clipboard imports. **rio** supports a few file formats not imported by **SchemaOnRead** such as fixed format files, FORTRAN data files, and clipboard imports. **SchemaOnRead** similary offers several formats not supported by **rio** such as deep XML, BMP, JPEG, and PNG files. Unlike **SchemaOnRead**, **rio** includes functions for writing as well as reading. Unlike **rio**, **SchemaOnRead** includes functions for recursively reading directories and offers an interface that is easily extensible by end users.

The **foreign** package (R Core Team et al., 2015) provides functions for reading a range of file types including Weka, Epi Info, SPSS, Stata, Systat files. **SchemaOnRead** uses **foreign** for reading these types of files. Unlike **SchemaOnRead**, **foreign** uses different user function calls to select the format of the file being imported. Unlike **foreign**, **SchemaOnRead** provides recursive reading of folders, is designed to be easily extended by end users to new file formats, and checks file extensions to determine formats.

The **readbitmap** package (Jefferis, 2015) provides functions for reading BMP, JPEG, and PNG files. **SchemaOnRead** uses **readbitmap** for reading BMP, JPEG and PNG files. Unlike **SchemaOnRead**, **readbitmap** uses magic numbers rather than extensions to identify file formats [2]. Unlike **readbitmap**, **SchemaOnRead** provides recursive reading of folders and is designed to be easily extended by end users to new file formats.

Why "SchemaOnRead?"

Schema-on-read (Deutsch, 2013), (Mendelevitch, 2013), (Jacobsohn and Delurey, 2014) is an agile approach to data storage and retrieval that defers investments in data organization until production queries need to be run by working with data directly in native form. Schema-on-read functions have been implemented in a wide range of analytical systems including Hadoop (Hadoop Team, 2015), (Schau, 2015), Splunk (Bitincka et al., 2012), Apache Spark (Spark Team, 2015), Apache Flink (Markl, 2014), and even relational databases (Liu and Gawlick, 2015). It is also possible to use machine learning tools to extract schemas from source data (Yeh et al., 2013).

The R Package SchemaOnRead

The **SchemaOnRead** R package defines four public functions:

- `schemaOnRead(path = ".",processors = defaultProcessors(),verbose = FALSE)` processes the given path using the provided lsit of processors optinally printing its progress on the console.

- `defaultProcessors()` returns a complete list of built-in processors in the recommended execution order.

- `simpleProcessors()` returns a minimal list of built-in processors in the recommended execution order.

- `checkExtensions(path = ".",extensions = NULL)` returns true if the path exists and, if an extensions list is provided, the extension of the path is in extensions list.

The `schemaOnRead` function is used to read source material (e.g., files and folders).

The SchemaOnRead package uses a recursive implementation. The initial user function call, `schemaOnRead` iterates over the given list of processors, invoking each in turn until one returns a non-null value. Processors are sequentially invoked in the order given by the input list, scanning from index number one upwards. Processing continues as long as each processor returns null. The results from the first processor to return a non-null value is stored as the content for the entry and processing of that entry stops. All of the results are stored in a named list. The order of the resulting list is the

---

[2]Magic numbers (Wikipedia, 2015) are special values in files that represent the file format. Magic numbers are commonly stored as special values encoded in file headers and footers. The first two bytes of JPEG files in hexidecmial are FF and D8 and the last two bytes are FF and D9. The first six bytes of GIF files in hexidecimal are 47, 49, 46, 38, 37, and 61 (GIF87a in ASCII) or 47, 49, 46, 38, 39, and 61 (GIF89a in ASCII). The first eight bytes of PNG files in hexidecmial are 89, 50, 4E, 47, 0D, 0A, 1A, and 0A) which, in part, spells PNG in ASCII.

order given by the file system. The variable names are taken from the entry names (e.g., file or folder names). Files or folders with names that do not conform to standard R variable naming requirements can be accessed using single quote notation (e.g., results$'Nonconforming Name').

An example processor for Microsoft Excel spreadsheets is shown in Figure 9. In this example, the entry identified by the path string is checked to see if it exists as a file. If it does, then the file name is extension is checked. If it matches then the processor attempts to read the file.

```r
## Define the XLS and XLSX spreadsheet file processor.
schemaOnReadProcessXLSandXLSXFile <- function(path = ".",
        processors = schemaOnReadDefaultProcessors(), verbose = FALSE) {

        ## Check the given path.
        if ((file.exists(path)) &&
                ((tolower(tools::file_ext(path)) == "xls") ||
                 (tolower(tools::file_ext(path)) == "xlsx"))) {

                ## Create the results holder.
                results <- list()

                ## Attempt to read the file.
                workbook <- XLConnect::loadWorkbook(path)

                ## Scan the worksheets.
                for (worksheet in XLConnect::getSheets(workbook)) {

                        ## Define the variable name.
                        variable <- gsub("[^[:alnum:].]", "_", worksheet)
                        while (eval(parse(
                                text = paste("exists(\"results$", variable, "\")",
                                        sep = "")))) {
                                variable <- paste(variable, "_A", sep = "")
                        }

                        ## Setup the processing command.
                        command <- paste("results$", variable,
                                " <- XLConnect::readWorksheet",
                                "(workbook, sheet = worksheet)", sep = "")

                        ## Evaluate the processing command.
                        eval(parse(text = command))

                }

                ## Return the results.
                return(results)

        } else {

                ## Return the default value.
                return(NULL)

        }

}
```

Figure 9: An Example Processor for Microsoft Excel XLS and XLSX Spreadsheets

The main goal of a processor is to read each acceptable entry into R in an easily usable format. Examples include the production of lists and data frames. The main output of SchemaOnRead is thus intended to be a nested tree of lists, with data frames in some of the leaves the tree. The example in Figure 1 does this by scanning the worksheets in a given workbook and converting each into a data frame. The result is a list of data frames with each data frame entry identified using the name of the corresponding worksheet. Note that the worksheet names are checked to insure that they correspond to valid R variable names for convenient user access.

The postconditions for each processor are that the processor or one of its descendants either successfully processes the entry and returns a non-null result or fail to process the entry and return null. If the entry is successfully processed then SchemaOnRead will perform no further processing on the item. If the item was not successfully processed then SchemaOnRead will use its remaining processors list to attempt to process the entry.

Several special processors are defined for SchemaOnRead. These include processors for nonexistent entries, directories, and entries of unknown types.

The `schemaOnReadProcessEntryDoesNotExist` processor returns null if the given entry exists and returns the value `"Entry Does Not Exist"` if not. It is meant to be the first processor in most lists to intercept nonexistent entries before they waste execution time in other processors. Occasionally, special processing may needed for nonexistent entries so these processors should run first.

The `schemaOnReadProcessDirectory` processor handles directories as previously discussed. It is intended to be the second processor to run in normal lists.

The `schemaOnReadProcessDefaultFile` processor accepts all entries that exist and returns the `"File Type Unknown"` string. It normally runs last to insure a value for unrecognized file types.

SchemaOnRead includes predefined two processing lists. The default processing list is used for SchemaOnRead entry processing. The simple processing list provides an easy starting point for user-defined processor lists.

Twenty-one unit tests are defined for the SchemaOnRead package. These tests are implemented using the **testthat** R package (Wickham, 2015). The current version of SchemaOnRead passes all of the defined tests.

Summary

As we have discussed, schema-on-read is a powerful new option for data storage and retrieval. Schema-on-read functions have been implemented in a wide range of analytical systems, most notably Hadoop. SchemaOnRead uses R's flexible data representations to provide transparent and convenient support for the schema-on-read paradigm in R. This paper's contribution is that it introduces and describes the **SchemaOnRead** package and compares it to related R packages.

Acknowledgements

Bibliography

L. Bitincka, A. Ganapathi, and S. Zhang. Experiences with workload management in splunk. In *Workshop on Management of Big Data Systems*, pages 25–30, 2012. [p4]

C. H. Chan, G. C. H. Chan, T. J. Leeper, and C. Gandrud. CRAN rio Package, Version 0.2. https://cran.r-project.org/web/packages/rio/index.html, 2015. [p4]

T. Deutsch. Why is schema on read so useful? http://www.ibmbigdatahub.com/blog/why-schema-read-so-useful, 2013. [p4]

Hadoop Team. Apache hadoop. http://hadoop.apache.org, 2015. [p4]

M. Jacobsohn and M. Delurey. How the data lake works. https://www.boozallen.com/content/dam/boozallen/documents/Data_Lake.pdf, 2014. [p4]

G. Jefferis. CRAN readbitmap Package, Version 0.1-4. https://cran.r-project.org/web/packages/readbitmap/index.html, 2015. [p4]

Z. H. Liu and D. Gawlick. Management of flexible schema data in rdbmss - opportunities and limitations for nosql. In 7^{th} *Biennial Conference on Innovative Data Systems Research*, 2015. [p4]

V. Markl. Breaking the chains: On declarative data analysis and data independence in the big data era. In *Proceedings of the VLDB Endowment*, volume 7, pages 1730–1733, 2014. [p4]

O. Mendelevitch. Apache hadoop and data agility. http://hortonworks.com/blog/hadoop-data-agility/, 2013. [p4]

R Core Team, R. Bivand, V. J. Carey, S. DebRoy, S. Eglen, R. Guha, N. Lewin-Koh, M. Myatt, B. Pfaff, B. Quistorff, F. Warmerdam, S. Weigand, and Free Software Foundation, Inc. CRAN foreign Package, Version 0.8-66. https://cran.r-project.org/web/packages/foreign/index.html, 2015. [p4]

RStudio. RStudio. https://www.rstudio.com, 2015. [p1]

A. Schau. Schema-on-read in action. http://blog.cask.co/2015/03/schema-on-read-in-action/, 2015. [p4]

Spark Team. Apache spark. http://spark.apache.org, 2015. [p4]

H. Wickham. CRAN testthat Package, Version 0.10.0. https://cran.r-project.org/web/packages/testthat/index.html, 2015. [p6]

Wikipedia. Magic number (programming). http://en.wikipedia.org/wiki/Magic_number_(programming), 2015. [p4]

E. Yeh, J. Niekrasz, and D. Freitag. Unsupervised discovery and extraction of semi-structured regions in text via self-information. In *Proceedings of the 2013 Workshop on Automated Knowledge Base Construction*, pages 103–107, 2013. [p4]

Michael J. North
Argonne National Laboratory
9700 South Cass Avenue
Argonne, IL 60439 USA
north@anl.gov