

anchoredDistr: a Package for the Bayesian Inversion of Geostatistical Parameters with Multi-type and Multi-scale Data

by Heather Savoy, Falk Heße, and Yoram Rubin

Abstract The Method of Anchored Distributions (MAD) is a method for Bayesian inversion designed for inferring both local (e.g. point values) and global properties (e.g. mean and variogram parameters) of spatially heterogeneous fields using multi-type and multi-scale data. Software implementations of MAD exist in C++ and C# to import data, execute an ensemble of forward model simulations, and perform basic post-processing of calculating likelihood and posterior distributions for a given application. This article describes the R package **anchoredDistr** that has been built to provide an R-based environment for this method. In particular, **anchoredDistr** provides a range of post-processing capabilities for MAD software by taking advantage of the statistical capabilities and wide use of the R language. Two examples from stochastic hydrogeology are provided to highlight the features of the package for MAD applications in inferring anchored distributions of local parameters (e.g. point values of transmissivity) as well as global parameters (e.g. the mean of the spatial random function for hydraulic conductivity).

Introduction

The field of Geostatistics originated in the 1950s with the pioneering work of Krige (1951) and Matheron (1962) who tried to estimate the characteristics of subsurface properties with the limited measurements typically available in this field. This scarcity, caused by the high exploration costs, is exacerbated by the strong heterogeneity that many such subsurface properties exhibit. Both these factors combined make it impossible to describe any subsurface process with certainty, therefore necessitating the application of statistical tools. Today, Geostatistics is used in many fields of Earth Science such as Geology (Hohn, 1962), Hydrogeology (Kitanidis, 2008), plus Hydrology and Soil science (Goovaerts, 1999). To meet this demand, many software packages have been developed that provide practitioners and scientists alike with the much needed tools to apply Geostatistics. In R, the best collection of such tools is arguably found in the **gstat** package (Pebesma, 2004) developed and maintained by Pebesma and colleagues. With **gstat**, it is possible to estimate (Kriging) and simulate (Gaussian process generation) heterogeneous fields in one, two or three dimensions, therefore providing necessary tools for geostatistical analysis.

Any such statistical analysis should draw on all available data that are connected to the variable of interest to infer, i.e. to learn about, its spatial distribution as much as possible. Examples for such spatially distributed variables in Earth Sciences would be, e.g. the hydraulic conductivity of an aquifer, evapotranspiration rates of different land surface areas, and soil moisture. In classical statistics, such information may consist of measurements of the variable itself or so-called local variables. Here, local means that a point-by-point relationship between both variables exists. However, many data are non-local, which means they are connected to the variable of interest via a complicated forward model. For instance, hydraulic conductivity may be connected by a solute transport model to break-through curves of said solutes and soil moisture may be connected by a hydraulic catchment model to river discharge. To learn about the input from the output of such forward models means to invert them, hence the name inversion for such techniques.

The Method of Anchored Distributions (MAD) provides a Bayesian framework for the geostatistical inversion of spatially heterogeneous variables. MAD solves the aforementioned problem by converting non-local data into equivalent local data using the tools of Bayesian inference. The result of such a conversion is the consistent representation of all data (local and non-local) as local data only, which is then amendable to further geostatistical analysis (Rubin et al., 2010). So far, applications of MAD have been focused on Hydrogeology (Murakami et al., 2010; Chen et al., 2012; Heße et al., 2015) as well as Soil Science (Over et al., 2015). However, given the explanations above, MAD is in no way limited to these fields and can be employed wherever non-local data need to be incorporated into a geostatistical framework. This generality also extends to the spatial model being inferred. While there are R packages utilizing Bayesian inference for spatial models such as **spBayes** (Finley et al., 2015), **spTimer** (Bakar and Sahu, 2015), and **INLA** (Lindgren and Rue, 2015, software available from <http://www.r-inla.org/>), these packages have several constraints compared to **anchoredDistr**. First, each method assumes a Gaussian process for the spatial variability. MAD has no inherent distributional assumptions, which allows its application to a wide variety of scenarios where, for

example, Gaussian fields are not justified. In addition, these packages are either geared toward large data sets (**spBayes** and **spTimer**) or applied to only local data (**spBayes**, **spTimer**, and **INLA**) while MAD focuses on addressing uncertainty due to sparse data sets by incorporating non-local data. Finally, MAD employs a non-parameteric likelihood estimation, which allows for great flexibility, in particular for non-linear forward models. The presented R package **anchoredDistr** provides an interface to the C# implementation of MAD. It allows post-processing of calculating likelihood and posterior distributions as well as visualization of the data.

The Method of Anchored Distributions

Equation 1 displays the general procedure of Bayesian inference where θ represents the parameters of the variable being inferred (e.g. hydraulic conductivity) and z represents the data informing the inference:

$$p(\theta|z) \propto p(\theta) p(z|\theta). \quad (1)$$

An important element of MAD is a strict classification of all data into local z_a and non-local data z_b , with the latter being the target of inversion. MAD employs Bayesian inference in the realm of Geostatistics by expanding the supported parameters into θ for global parameters (describing overall trend and spatial correlation) and ϑ for local parameters. Since MAD is a Bayesian scheme, these θ and ϑ both have probability distributions. As mentioned above, MAD turns non-local data into equivalent local data ϑ by inverting the forward model that connects both. The non-local data therefore become anchored in space, hence the name Method of Anchored Distributions. Equation 2 displays the general form of MAD:

$$p(\theta, \vartheta|z_a, z_b) \propto p(\theta) p(\vartheta|\theta, z_a) p(z_b|\theta, \vartheta, z_a). \quad (2)$$

Open-source software implementations for applying the entirety of MAD are available both with a graphical interface and a command-line interface to guide users through connecting their forward models and random field generators and to execute the ensemble of forward simulations (a. Osorio-Murillo et al., 2015). This software (available at <http://mad.codeplex.com>) was inspired by the claim that inverse modeling will be widely applied in Hydrogeology only if user-friendly software tools are available (Carrera et al., 2005).

The package **anchoredDistr** described here focuses on extending the post-processing capabilities of MAD software, particularly the calculation of the likelihood distribution $p(z_b|\theta, \vartheta, z_a)$ and the posterior distribution $p(\theta, \vartheta|z_b, z_a)$ after the ensemble of forward model simulations is already complete. The MAD# software has basic post-processing capabilities, but does not offer the degree of flexibility as R for the post-processing analysis. For example, when handling z_b in the form of time series, dimension reduction techniques are necessary for calculating the likelihood values. By having the R package **anchoredDistr**, users have the support to attach whichever applicable technique for their data.

General workflow

In the current version of **anchoredDistr**, which only handles the post-processing of a MAD application, it is assumed that prior distributions of local and global parameters, $p(\vartheta|\theta, z_a)$ and $p(\theta)$ respectively, have already been defined and sampled and that forward model simulations based on those samples have been executed either within the MAD# software or by other means of batch execution. If the MAD# software is used, this data is stored by MAD# in databases (extensions .xresult for project metadata and .xdata for each sample). The package **anchoredDistr** primarily consists of methods for the S4 class "MADproject" that extract and analyze data from these databases, i.e. handling information regarding the samples from the prior distributions and the resulting ensemble of simulated z_b data. If MAD# is not used, the information can be formatted into a "MADproject" manually. The usage of **anchoredDistr** will generally follow the workflow below (also see Figure 1):

1. Create "MADproject" object with new function (passing slot information if manually filling data)
2. Read data from MAD# databases, if being used, into "MADproject" object with readMAD
3. View the observations and realizations with plotMAD
4. Apply any necessary dimension reduction techniques to z_b with reduceData
5. Test the convergence of the likelihood distribution with respect to the number of realizations with testConvergence (return to MAD software to run additional realizations if unsatisfactory)
6. Calculate likelihood and posterior distributions with calcLikelihood and calcPosterior, respectively

7. View the posterior distribution with plotMAD.

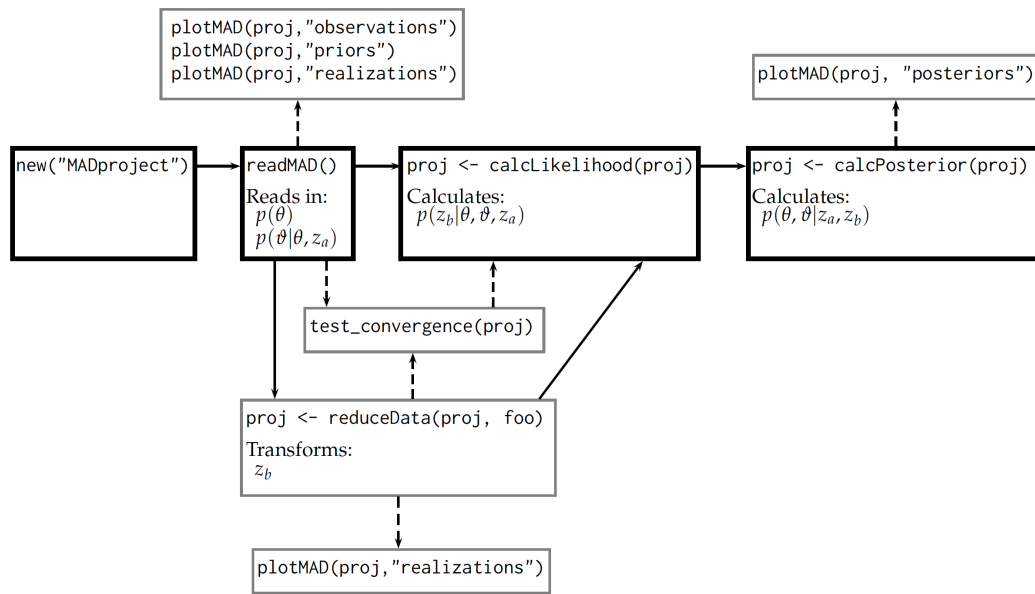


Figure 1: Schematic of utilizing **anchoredDistr** for MAD post-processing if the MAD# is used. Solid arrow lines indicate the fundamental workflow while dashed arrow lines are optional.

To install the **anchoredDistr** package, the release version is available from CRAN:

```
install.packages("anchoredDistr")
library(anchoredDistr)
```

Alternatively, the development version can be obtained by using the **devtools** package (Wickham and Chang, 2016) to download the necessary files from GitHub:

```
library(devtools)
install_github("hsavoy/anchoredDistr")
library(anchoredDistr)
```

Other packages used by **anchoredDistr** include **RSQLite** (Wickham et al., 2014) for reading from MAD databases, **np** (Hayfield and Racine, 2008) for estimating non-parametric density distributions, **plyr** (Wickham, 2011) and **dplyr** (Wickham and Francois, 2016) for efficient data manipulation, and **ggplot2** (Wickham, 2009) for plotting. The methods included in **anchoredDistr** are listed in Table 1 and two examples utilizing these methods are provided next.

Method	Description
<code>readMAD()</code>	Reads data from databases generated by MAD software
<code>reduceData()</code>	Applies dimension reduction to z_b time series
<code>testConvergence()</code>	Tests for convergence of likelihood values for increasing number of realizations
<code>calcLikelihood()</code>	Calculates the likelihood values for the samples
<code>calcPosterior()</code>	Calculates the posterior values for the samples
<code>plotMAD()</code>	Plots the observations, realizations, reduced data, and/or posteriors

Table 1: The methods for the "MADproject" S4 class provided by **anchoredDistr**.

Example 1: aquifer characterization with steady-state hydraulic head from multiple wells

Scenario setup

In this first example, we will use the tutorial example available from the MAD website <http://mad.codeplex.com>. Within the **anchoredDistr** package, this tutorial example is available as MAD# databases, as well as a "MADproject" object accessed by `data(tutorial)`. The variable of interest is transmissivity T , an aquifer property that represents how much water can be transmitted horizontally through an aquifer. We will use the one-dimensional heterogenous field of the decimal log transform of T (see Figure 2) as our baseline field from which we can generate virtual measurements and validate our resulting posterior distributions. The field was generated as a Gaussian process by the **gstat** package in R with a mean $\mu_{\log_{10} T} = -2$ and an exponential covariance function with a variance $\sigma_{\log_{10} T}^2 = 0.4$ and length scale $l_{\log_{10} T} = 3$ m. Within the scope of this example, we assume these global parameter values to be known. Furthermore, we assume that we have local data in the form of measurements of T at three different locations. In addition, non-local data are available in the form of head measurements (indication of water pressure) at the same locations. The forward model used to solve the groundwater flow equation and relate T to head is the software MODFLOW-96 (Harbaugh and McDonald, 1996), part of the open source MODFLOW series that is the industry standard for groundwater modeling. To convert the non-local data into equivalent local data of T , we will place four anchors at selected unmeasured locations. The number of anchors needs to be justified by the data content of the measurements such that the complexity of the model does not become disproportionate to the information available. The locations of these anchors reflect locations where there is no other local data available but there is non-local data nearby for conversion (see Yang et al. (2012) for more discussion on anchor placement). The locations of the measurements and anchors are depicted in Figure 2. The prior distributions for these anchors are based on simple kriging with the local data z_a for conditioning and the known Gaussian process for the covariance function:

$$p(\theta_i | \theta, z_a) = \mathcal{N}(\mu = \hat{Z}(y_i), \sigma^2 = \text{Var}(Z(y_i) - \hat{Z}(y_i))), \quad (3)$$

where Z generally represents $\log_{10} T$, y_i is the y -coordinate of the i^{th} anchor, $\hat{Z}(y_i)$ is the kriging estimate at the i^{th} anchor, and $\text{Var}(Z(y_i) - \hat{Z}(y_i))$ is the kriging variance at the i^{th} anchor. The goal of the example is to compare the posterior distributions of the four anchors resulting from the inversion to their prior distributions which will indicate the information gain from the inclusion of the non-local data z_b .




Figure 2: The one-dimensional baseline field of $\log_{10} T$ used in Example 1 with locations of measurements (co-located z_a and z_b) marked along with the anchors to be inferred.

Reading and viewing data

In the first step, a "MADproject" object is created with the `new()` function. Three arguments must be provided to read the MAD databases: `madname` (the name of the MAD project, e.g. the filename for the .xmad database), `resultname` (the name of the result from MAD, e.g. the result folder name), and `xpath` (the path to where the .xresult database and result folder are located). These three arguments ensure the MAD databases can be read by the method `readMAD`, which will read in the prior distribution samples for the global and local parameters plus the observations and forward model predictions for the z_b . Note that **anchoredDistr** could be used independently of the MAD software, if desired, as long

as the slots filled in by readMAD (see Table 2) are provided manually (see next example). To create a "MADproject" object for this tutorial example, the code below will read the MAD# databases stored in the **anchoredDistr** package files.

```
tutorial <- new("MADproject", madname="Tutorial", resultname="example1",
  xpath=paste0(system.file("extdata", package = "anchoredDistr"),"//")
tutorial <- readMAD(tutorial, 1:3)
```

Slot	Description	Source
madname	MAD project name	user provided
resultname	MAD result name	user provided
xpath	Path to .xresult database	user provided
numLocations	Number of z_b locations	readMAD
numTimesteps	Number of time steps measured at each z_b locations	readMAD
numSamples	Number of samples drawn from prior distributions	readMAD
numAnchors	Number of local parameters / anchors placed in field	readMAD
numTheta	Number of random global parameters to infer	readMAD
truevalues	True values for the parameters to infer, if known	readMAD
observations	Observed values of the z_b locations and time steps	readMAD
realizations	Simulated values of the z_b locations and time steps	readMAD
priors	Samples from the prior distributions of each parameter	readMAD
likelihoods	Likelihood values for each sample	calcLikelihoods
posteriors	Posterior values for each sample of each parameter	calcPosteriors

Table 2: The slots for the "MADproject" S4 class provided by **anchoredDistr**.

The prior distributions can be viewed by calling the plotMAD function with the "MADproject" object and the string "priors" (see below). Figure 3 shows the prior distributions for the four anchors in Example 1. The distributions roughly follow a Gaussian distribution due to the baseline field being a Gaussian field and the prior distributions based on the kriging mean and variance at these four locations from the z_a data and the known spatial random function. The x-axis labels are pulled from the "MADproject" object's priors slot, which contains the random parameter names as provided in the MAD software.

```
plotMAD(tutorial, "priors")
```

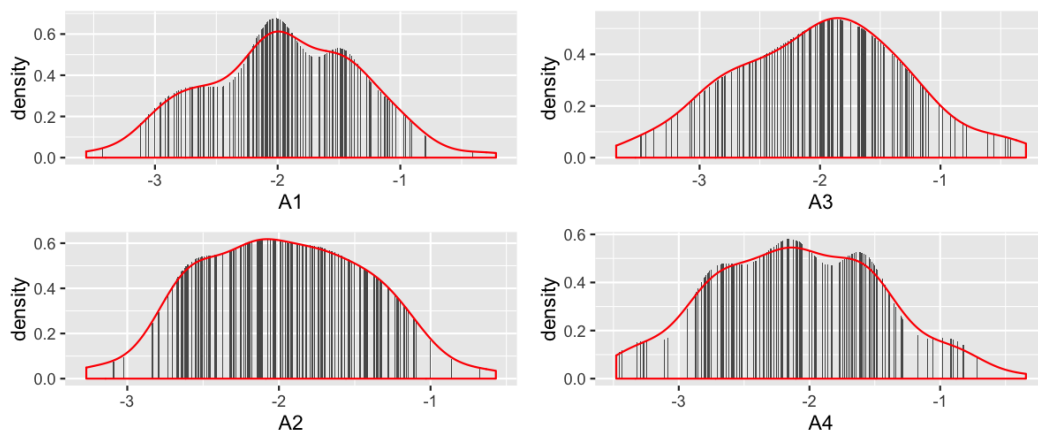


Figure 3: The relative frequency (gray bars) and estimated density (red line) of the prior distributions for the four anchor locations based on samples supplied in Example 1.

Calculating likelihoods and posteriors

After the information contained in the MAD databases has been read into the "MADproject" object, the likelihood and posterior distributions can be calculated by calcLikelihood and calcPosterior, respectively. The method calcLikelihood uses non-parametric kernel density estimation (from the

package **np**) to estimate the probability density of measured inversion data from the probability density function of inversion data simulated from the realizations per sample. The method `calcPosterior` multiplies the resulting likelihood distribution across the samples and the provided prior distribution to calculate the posterior.

First, we can call the `testConvergence()` method to visually inspect if we have enough realizations for the likelihood values of samples to converge (this method calls the `calcLikelihood` internally to perform this test). Figure 4 depicts this qualitative convergence test for Example 1 by plotting the likelihood values of a sample with increasing number of realizations. In order to prevent cluttering, the default number of samples to display is set to seven samples randomly selected from those available in the project. Convergence is achieved when the likelihood stabilizes with increasing realizations. For this example, it appears that the log likelihood of the samples have started to stabilize by 50 realizations, but more realizations may be warranted.

The posterior distributions for each random parameter can be seen by calling `plotMAD` with the "MADproject" object and the string "posteriors". Figure 5 shows the posteriors for Example 1 along with the prior distribution and the true values for each of the four anchors. The posterior distributions for Anchors 2 and 3, which were surrounded by z_b measurements, show an increase in probability near the true value, indicating a successful information transfer from the non-local z_b into equivalent local data.

```
testConvergence(tutorial)
tutorial <- calcLikelihood(tutorial)
tutorial <- calcPosterior(tutorial)
plotMAD(tutorial, "posteriors")
```

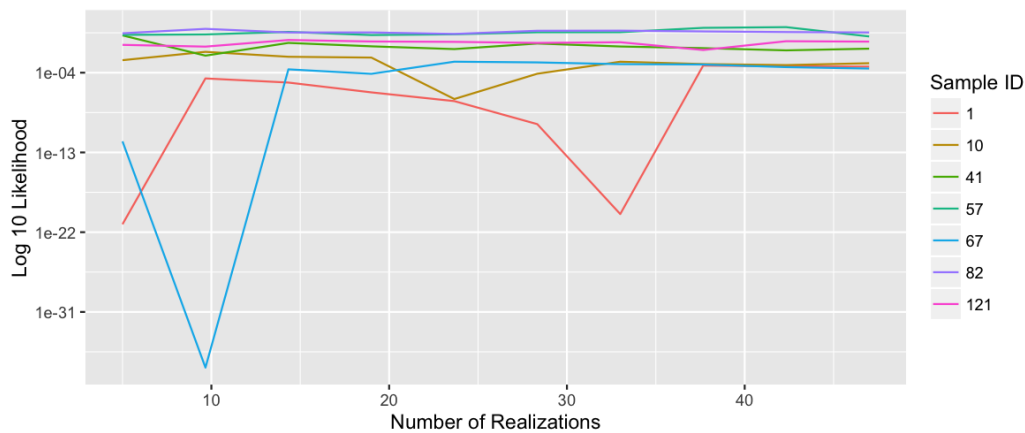


Figure 4: Convergence testing for Example 1 by plotting the decimal log of likelihood of a collection of randomly selected samples with increasing number of realizations.

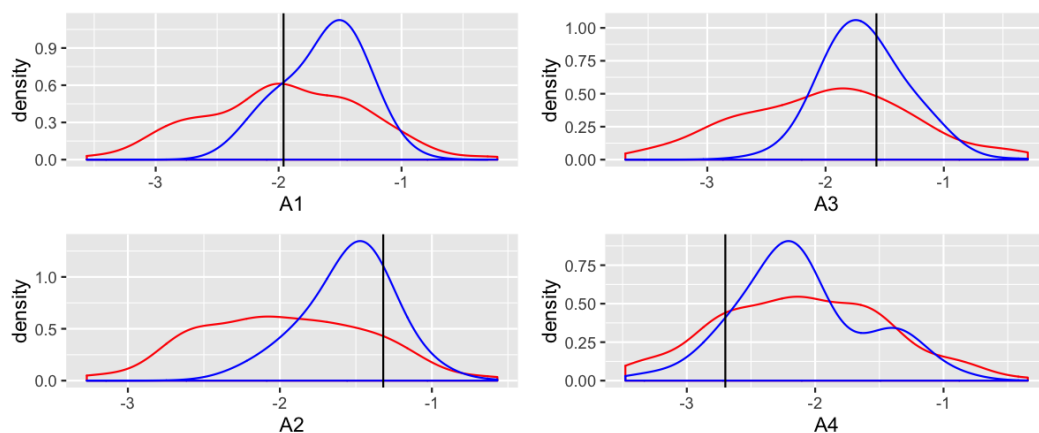


Figure 5: The prior (red) and posterior (blue) distributions with the true value (black) for the four anchor locations in Example 1.

Example 2: aquifer characterization with one pumping drawdown curve

Scenario setup

The second example depicts a different aquifer characterization scenario for a two-dimensional field where the natural log transform of hydraulic conductivity (K) is assumed to be an isotropic Gaussian field with variance $\sigma_{\ln K}^2 = 1$ and length scale $l_{\ln K} = 10\text{m}$ but unknown mean $\mu_{\ln K}$ (Figure 6). There are no anchors placed in this example, leaving the mean as the only parameter to infer. Unlike Example 1, Example 2 is therefore a demonstration of how MAD can be employed as a regular Bayesian inversion scheme, too. The prior distribution for global parameters ideally come from previous knowledge of similar sites, e.g. the distribution of mean $\ln K$ observed at other aquifers with the same geological setting. For this example, we will compare three equally spaced samples for $\ln K$ to represent a uniform prior distribution for the mean. The data include four local data z_a (K) at four different locations and one non-local data series z_b (hydraulic head drawdown) at a single location (see Figure 6). The z_b location provides 100 time steps, i.e. data points, of drawdown measurements (Figure 7). The forward model used to solve the groundwater flow equation and relate K to drawdown is OpenGeoSys (Kolditz et al., 2012), an open source software that simulates a variety of subsurface processes. This second example uses a different forward model than the first example to showcase the MAD software's modular design that does not assume or rely on specific forward models. The observation, realizations, and prior sample data for this example is provided within the package as external data that can be created with new as shown below, as well as a pre-made "MADproject" object accessed with data(pumping).

```
load(system.file("extdata", "pumpingInput.RData", package = "anchoredDistr"))
pumping <- new("MADproject",
  numLocations = 1,
  numTimesteps = 100,
  numSamples = 50,
  numAnchors = 0,
  numTheta = 1,
  observations = obs,
  realizations = realizations,
  priors = priors)
```




Figure 6: The two-dimensional baseline field of $\ln K$ used in Example 2 with the location of measurements marked.

When the pumping dataset is initially loaded, we can view the observation of z_b , i.e. drawdown time series (Figure 7), the prior distribution of the three samples (Figure 8), and the interquartile range of the time series simulated by the forward model for the samples (Figure 9).

```
plotMAD(pumping, "observations")
```

```
plotMAD(pumping, "priors")
plotMAD(pumping, "realizations")
```




Figure 7: The observed time series of hydraulic head drawdown to be used as non-local data z_b in Example 2.




Figure 8: The histogram (gray bars) and estimated density (red line) of the prior distributions for the mean $\ln K$ Example 2.

Applying dimension reduction to time series

Even though we have the time series of drawdown, we cannot use these 100 individual values to calculate the likelihood because they are correlated and the multivariate likelihood distribution would be 100-dimensional. Such dimensionality would require an unrealistic number of realizations to resolve, known as "the curse of dimensionality." To overcome this obstacle, dimension reduction is needed and the method to use depends on the type of non-local data z_b . For this example, we will simply use the `min` function to collect the minimum head value in the time series since the observed head reduces and converges to a stable head value with time (Figure 7). The **anchoredDistr** package can handle any non-parameterized function, such as `min`, or a parameterized function if initial values for each parameter are given and the `nls` function (R Core Team, 2016) can perform the fitting (see the package vignette for an example). The `reduceData` function is used to perform the dimension reduction on the time series:

```
pumping.min <- reduceData(pumping, min)
plotMAD(pumping.min, "realizations")
```

The `reduceData` function returns a "MADproject" object with a `realizations` slot with reduced dimensions. The reduced data can be viewed by calling `plotMAD` with the string "realizations". The plot shows the distributions of each parameter for each sample. In this case, Figure 10 shows the




Figure 9: The observed time series of drawdown at the z_b location along with the inter-quartile range of simulated values for each time step for the three samples.




Figure 10: The reduced z_b data (minimum of drawdown curve) for Example 2. Distributions are estimated from the realizations' reduced data per sample.




Figure 11: The prior (red) and posterior (blue) distributions with the true value (black) for the mean $\ln K$ locations in Example 2.

minimum head value distribution for the three samples, which will be used to calculate the three likelihood samples.

With this new "MADproject" object, `calcLikelihoods` and `calcPosteriors` can be called. In Figure 11, the posterior distributions are shown for the three samples along with the true value of -10 . The posterior distribution assigns greater probability toward the true value.

```
pumping.min <- calcLikelihood(pumping.min)
pumping.min <- calcPosterior(pumping.min)
plotMAD(pumping.min, "posteriors")
```

Summary

The examples given above show how the **anchoredDistr** package allows flexible post-processing of results by virtue of the MAD software such that users can apply their own post-processing analyses, such as dimension-reduction techniques. The first example shown here is available as external and internal datasets in the **anchoredDistr** package. The second example is also included in **anchoredDistr** and is further detailed in the package vignette. The release version of the **anchoredDistr** package is hosted on CRAN and the development version is hosted on GitHub, which can be accessed by calling `devtools::install_github("hsavoy/anchoredDistr")` or by downloading from <http://hsavoy.github.io/anchoredDistr>.

Acknowledgements

This work was supported by the National Science Foundation under grant EAR-1011336, "The Method of Anchored Distributions (MAD): Principles and Implementation as a Community Resource." Any opinions, findings and conclusions or recommendations expressed in this material are those of the authors and do not necessarily reflect the views of the National Science Foundation.

Bibliography

- C. a. Osorio-Murillo, M. W. Over, H. Savoy, D. P. Ames, and Y. Rubin. Software Framework for Inverse Modeling and Uncertainty Characterization. *Environmental Modelling & Software*, 66:98–109, 2015. ISSN 13648152. URL <https://doi.org/10.1016/j.envsoft.2015.01.002>. [p2]
- K. S. Bakar and S. K. Sahu. spTimer: Spatio-Temporal Bayesian Modeling Using R. *Journal of Statistical Software*, 63(15):32, 2015. ISSN 1548-7660. [p1]
- J. Carrera, A. Alcolea, A. Medina, J. Hidalgo, and L. J. Slooten. Inverse Problem in Hydrogeology. *Hydrogeology Journal*, 13(1):206–222, 2005. ISSN 14312174. URL <https://doi.org/10.1007/s10040-004-0404-7>. [p2]
- X. Chen, H. Murakami, M. S. Hahn, G. E. Hammond, M. L. Rockhold, J. M. Zachara, and Y. Rubin. Three-Dimensional Bayesian Geostatistical Aquifer Characterization at the Hanford 300 Area Using Tracer Test Data. *Water Resources Research*, 48(6):W06501, 2012. ISSN 0043-1397. URL <https://doi.org/10.1029/2011wr010675>. [p1]
- A. O. Finley, S. Banerjee, and A. E. Gelfand. spBayes for Large Univariate and Multivariate Point-Referenced Spatio-Temporal Data Models. *Journal of Statistical Software*, 63(13):1–28, 2015. ISSN 1548-7660. URL <http://www.jstatsoft.org/v63/i13>. [p1]
- P. Goovaerts. Geostatistics in soil science: State-of-the-art and perspectives. *Geoderma*, 89(1–2):1 – 45, 1999. ISSN 0016-7061. URL [https://doi.org/10.1016/s0016-7061\(98\)00078-0](https://doi.org/10.1016/s0016-7061(98)00078-0). [p1]
- W. Harbaugh and M. G. McDonald. User's Documentation for MODFLOW-96, an Update to the U.S. Geological Survey Modular Finite-Difference Ground-Water Flow Model, Open File Report 96-485. Technical report, U.S. Geological Survey, 1996. [p4]
- T. Hayfield and J. S. Racine. Nonparametric econometrics: The np package. *Journal of Statistical Software*, 27(5), 2008. URL <http://www.jstatsoft.org/v27/i05/>. [p3]
- F. Heße, H. Savoy, C. A. Osorio-Murillo, J. Sege, S. Attinger, and Y. Rubin. Characterizing the Impact of Roughness and Connectivity Features of Aquifer Conductivity Using Bayesian Inversion. *Journal of Hydrology*, 531:73–87, 2015. ISSN 00221694. URL <https://doi.org/10.1016/j.jhydrol.2015.09.067>. [p1]

- M. Hohn. *Geostatistics and Petroleum Geology (2nd Ed.)*. Kluwer, 1962. [p1]
- P. Kitanidis. *Introduction to Geostatistics: Applications in Hydrogeology*. Cambridge University Press, 2008. [p1]
- O. Kolditz, S. Bauer, L. Bilke, N. Böttcher, J. O. Delfs, T. Fischer, U. J. Görke, T. Kalbacher, G. Kosakowski, C. I. McDermott, C. H. Park, F. Radu, K. Rink, H. Shao, H. B. Shao, F. Sun, Y. Y. Sun, A. K. Singh, J. Taron, M. Walther, W. Wang, N. Watanabe, Y. Wu, M. Xie, W. Xu, and B. Zehner. OpenGeoSys: An Open-Source Initiative for Numerical Simulation of Thermo-Hydro-Mechanical/Chemical (THM/C) Processes in Porous Media. *Environmental Earth Sciences*, 67(2):589–599, 2012. ISSN 1866-6280. URL <https://doi.org/10.1007/s12665-012-1546-x>. [p7]
- D. G. Krige. A Statistical Approach to Some Basic Mine Valuation Problems on the Witwatersrand. *Journal of the Chemical, Metallurgical and Mining Society of South Africa*, 52(6):119–139, 1951. URL <https://doi.org/10.2307/3006914>. [p1]
- F. Lindgren and H. Rue. Bayesian Spatial Modelling with `\pkgR-INLA`. *Journal of Statistical Software*, 63(19):1–25, 2015. ISSN 1548-7660. URL <https://doi.org/10.18637/jss.v063.i19>. [p1]
- G. Matheron. *Traité De Géostatistique Appliquée*, volume 14. Editions Technip, Paris, 1962. [p1]
- H. Murakami, X. Chen, M. S. Hahn, Y. Liu, M. L. Rockhold, V. R. Vermeul, J. M. Zachara, and Y. Rubin. Bayesian Approach for Three-Dimensional Aquifer Characterization at the Hanford 300 Area. *Hydrology and Earth System Sciences*, 14(10):1989–2001, 2010. ISSN 1607-7938. URL <https://doi.org/10.5194/hess-14-1989-2010>. [p1]
- M. W. Over, U. Wollschlaeger, C. a. Osorio-Murillo, and Rubin. Bayesian Inversion of Mualem-Van Genuchten Parameters in a Multilayer Soil Profile: A Data-Driven Assumption-Free Likelihood Function. *Water Resources Research*, 51(2):861–884, 2015. URL <https://doi.org/10.1002/2013wr014956>. received. [p1]
- E. J. Pebesma. Multivariable geostatistics in s: The gstat package. *Computers & Geosciences*, 30:683–691, 2004. [p1]
- R Core Team. *R: A Language and Environment for Statistical Computing*. R Foundation for Statistical Computing, Vienna, Austria, 2016. URL <https://www.R-project.org/>. [p8]
- Y. Rubin, X. Chen, H. Murakami, and M. Hahn. A Bayesian Approach for Inverse Modeling , Data Assimilation , and Conditional Simulation of Spatial Random Fields. *Water Resources Research*, 46 (October 2009):1–23, 2010. ISSN 00431397. URL <https://doi.org/10.1029/2009wr008799>. [p1]
- H. Wickham. *Ggplot2: Elegant Graphics for Data Analysis*. Springer-Verlag, 2009. ISBN 978-0-387-98140-6. URL <http://ggplot2.org>. [p3]
- H. Wickham. The split-apply-combine strategy for data analysis. *Journal of Statistical Software*, 40(1): 1–29, 2011. URL <http://www.jstatsoft.org/v40/i01/>. [p3]
- H. Wickham and W. Chang. *Devtools: Tools to Make Developing R Packages Easier*, 2016. URL <https://CRAN.R-project.org/package=devtools>. R package version 1.11.1. [p3]
- H. Wickham and R. Francois. *Dplyr: A Grammar of Data Manipulation*, 2016. URL <https://CRAN.R-project.org/package=dplyr>. R package version 0.5.0. [p3]
- H. Wickham, D. A. James, and S. Falcon. *RSQLite: SQLite Interface for R*, 2014. URL <https://CRAN.R-project.org/package=RSQLite>. R package version 1.0.0. [p3]
- Y. Yang, M. Over, and Y. Rubin. Strategic Placement of Localization Devices (Such as Pilot Points and Anchors) in Inverse Modeling Schemes. *Water Resources Research*, 48(8):W08519, 2012. ISSN 00431397. URL <https://doi.org/10.1029/2012wr011864>. [p4]

Heather Savoy
Civil and Environmental Engineering
University of California, Berkeley
Berkeley, CA, USA
frystacka@berkeley.edu

Falk Hefße
Computational Hydrosystems

Helmholtz Centre for Environmental Research (UFZ)
Leipzig, Germany
falk.hesse@ufz.de

Yoram Rubin
Civil and Environmental Engineering
University of California, Berkeley
Berkeley, CA, USA
rubin@ce.berkeley.edu