

ManlyMix: An R Package for Manly Mixture Modeling

by Xuwen Zhu, Volodymyr Melnykov

Abstract Model-based clustering is a popular technique for grouping objects based on a finite mixture model. It has countless applications in different fields of study. The R package **ManlyMix** implements the Manly mixture model that allows modeling skewness within data groups and performs cluster analysis. **ManlyMix** is a powerful diagnostics tool that is capable of conducting investigation concerning the normality of variables upon fitting of a Manly forward or backward model. Theoretical foundations as well as description of functions are provided. All features of the package are illustrated with examples in great detail. The analysis of real-life datasets demonstrates the flexibility and usefulness of the package.

Introduction

Finite mixture models provide a powerful tool to model heterogeneous data. Their flexibility, close connection to cluster analysis, and interpretability make them increasingly appealing to researchers and practitioners these days. The applications of finite mixture modeling can be found in all fields, including medicine (Schlattmann, 2009), transportation (Park and Lord, 2009), dendrochronology (Michael and Melnykov, 2016), and environment science (Gillespie and Neale, 2006), just to name a few. The Bayes decision rule, applied to posterior probabilities obtained in the course of fitting a mixture model, yields clustering result. Such a procedure is called model-based clustering. It assumes the existence of a one-to-one correspondence between each distribution in the mixture model and underlying data group.

If all components in the model are Gaussian distributions, the mixture is called a Gaussian mixture model. Gaussian mixtures are very popular among practitioners due to their interpretability and simplicity. However, when there is severe skewness in data, Gaussian mixtures models do not provide a good fit to the data. As a result, model-based clustering might produce unsatisfactory results. In such cases, more flexible mixtures should be adopted. Some existing software packages that provide such functionality are listed in Table 1. Here, **mixmsn**, **EMMIXskew**, and **EMMIXuskew** packages are based on skew-normal and skew- t distributions, which are popular choices for modeling skewed data. On the other hand, **flowClust** is the only package that implements a transformation-based mixture model. It relies on the celebrated Box-Cox transformation to near-normality applied to all dimensions within the same mixture component. This leads to extra K parameters λ_k in the resulting mixture. The package is shown to model flow cytometry data effectively. In many applications, however, it is reasonable to assume that transformation parameters can vary not only from component to component but also from variable to variable. In this paper, we introduce the R package **ManlyMix** (Zhu and Melnykov, 2016b), which provides readers with an alternative approach to modeling and clustering skewed data. Manly mixture models (Zhu and Melnykov, 2016a) are constructed based on the Manly back-transformation applied to each variable in multivariate Gaussian components.

The **ManlyMix** package implements several functions associated with Manly mixture models including the core function for running the EM algorithm, the forward and backward model selection procedure for eliminating unnecessary transformation parameters, and the Manly K -means algorithm, which serves as an extension of the traditional K -means. Other capabilities of the package include computing a Manly mixture overlap, simulating datasets from a Manly mixture, constructing density or contour plots for a fitted model, and assessing the variability of estimated parameters. The highlights of **ManlyMix** include:

- providing an alternative approach to modeling heterogeneous skewed data;
- calling core functions from C for speed;

Package	Mixture components
flowClust (Lo et al., 2009)	t mixture with Box-Cox transformation
mixmsn (Prates et al., 2013)	scale skew-normal and skew- t
EMMIXskew (Wang et al., 2013)	restricted skew-normal and skew- t
EMMIXuskew (Lee and McLachlan, 2014)	unrestricted skew- t

Table 1: Existing R packages for mixture modeling of skewed data.

- providing excellent model interpretability through output of skewness parameters;
- preventing overfitting of the data by implementing model selection algorithms;
- offering effective assessment of mixture characteristics through the overlap calculation and variability assessment.

This paper is organized in the following way. A brief introduction to the Expectation-Maximization (EM) algorithm for Manly mixture models as well as the classification Expectation-Maximization (CEM) algorithm for Manly K -means is provided in the second section. In section "Package functionality and illustrative examples", a comprehensive description of all functions in **ManlyMix** is given along with the analysis of two real-life datasets. All features of the package are illustrated in great detail. Demo examples are constructed in section four for users to conduct further investigation of **ManlyMix**. In the last section, we provide a brief summary for the paper.

Methodological and algorithmic details

Manly mixture model

Consider a dataset X_1, \dots, X_n of size n , where X_i 's are p -variate independent observations that are identically distributed. The exponential (Manly) transformation to near normality is defined by

$$\mathcal{M}(X; \lambda) = \left(\frac{e^{\lambda_1 X_1} - 1}{\lambda_1}, \dots, \frac{e^{\lambda_p X_p} - 1}{\lambda_p} \right)^T,$$

where the distribution of $\mathcal{M}(X; \lambda_k)$ can be effectively approximated by multivariate normal distribution for an appropriate choice of λ (Manly, 1976). \mathcal{M}^{-1} represents the Manly back-transformation. This leads to a so-called Manly mixture model given by

$$g(x; \Psi) = \sum_{k=1}^K \tau_k \phi(\mathcal{M}(x; \lambda_k); \mu_k, \Sigma_k) \exp\{\lambda_k^T x\}, \quad (1)$$

where K is the number of components in the model, τ_k 's are mixing proportions such that $\sum_{k=1}^K \tau_k = 1$, and $\lambda_k = (\lambda_{k1}, \dots, \lambda_{kp})^T$ is a p -dimensional skewness vector which controls the transformation of the k th component. $\phi(\cdot; \mu_k, \Sigma_k)$ is the p -variate normal probability density function. μ_k and Σ_k are the mean vector and variance-covariance matrix of the k th component after transformation. Ψ , as the entire parameter vector, includes τ_k 's, μ_k 's and Σ_k 's.

To find the MLE of the parameter vector Ψ , the Expectation-Maximization (EM) algorithm (Dempster et al., 1977; McLachlan and Krishnan, 2008) needs to be employed. Each iteration of the EM algorithm consists of two steps, the E-step and M-step. Let s denote the iteration number. The E-step computes the posterior probabilities

$$\pi_{ik}^{(s)} = \frac{\tau_k^{(s-1)} \phi(\mathcal{M}(x_i; \lambda_k^{(s-1)}); \mu_k^{(s-1)}, \Sigma_k^{(s-1)}) \exp\{(\lambda_k^{(s-1)})^T x_i\}}{\sum_{k'=1}^K \tau_{k'}^{(s-1)} \phi(\mathcal{M}(x_i; \lambda_{k'}^{(s-1)}); \mu_{k'}^{(s-1)}, \Sigma_{k'}^{(s-1)}) \exp\{(\lambda_{k'}^{(s-1)})^T x_i\}} \quad (2)$$

based on the the parameter vector from the previous step, $\Psi^{(s-1)}$. The M-step updates the parameters in each iteration. The closed-form expressions are available for the parameters $\tau_k^{(s)}$, $\mu_k^{(s)}$, $\Sigma_k^{(s)}$ and are given by

$$\begin{aligned} \tau_k^{(s)} &= \frac{\sum_{i=1}^n \pi_{ik}^{(s)}}{n}, & \mu_k^{(s)} &= \frac{\sum_{i=1}^n \pi_{ik}^{(s)} \mathcal{M}(x_i; \lambda_k^{(s)})}{\sum_{i=1}^n \pi_{ik}^{(s)}}, \quad \text{and} \\ \Sigma_k^{(s)} &= \frac{\sum_{i=1}^n \pi_{ik}^{(s)} (\mathcal{M}(x_i; \lambda_k^{(s)}) - \mu_k^{(s)}) (\mathcal{M}(x_i; \lambda_k^{(s)}) - \mu_k^{(s)})^T}{\sum_{i=1}^n \pi_{ik}^{(s)}}. \end{aligned} \quad (3)$$

For λ_k , closed-form solution is not available and Nelder-Mead numerical optimization of the function

$$Q_k(\lambda_k | \Psi^{(s)})(\lambda_k) = \sum_{i=1}^n \pi_{ik}^{(s)} \left\{ \log \phi \left(\mathcal{M}(\mathbf{x}_i; \lambda_k); \sum_{i=1}^n \pi_{ik}^{(s)} \mathcal{M}(\mathbf{x}_i; \lambda_k) / \sum_{i=1}^n \pi_{ik}^{(s)}, \right. \right. \\ \left. \left. \sum_{i=1}^n \pi_{ik}^{(s)} \left(\mathcal{M}(\mathbf{x}_i; \lambda_k) - \frac{\sum_{i=1}^n \pi_{ik}^{(s)} \mathcal{M}(\mathbf{x}_i; \lambda_k)}{\sum_{i=1}^n \pi_{ik}^{(s)}} \right) \left(\mathcal{M}(\mathbf{x}_i; \lambda_k) - \frac{\sum_{i=1}^n \pi_{ik}^{(s)} \mathcal{M}(\mathbf{x}_i; \lambda_k)}{\sum_{i=1}^n \pi_{ik}^{(s)}} \right)^T \right) \right. \\ \left. + \lambda_k^T \mathbf{x}_i \right\} + \text{const} \quad (4)$$

gives us the estimated parameter vector.

The EM algorithm could be started with an initial partition of the data passed into the M-step. Or the E-step is run first with initial parameters $\tau_k^{(0)}, \mu_k^{(0)}, \Sigma_k^{(0)}, \lambda_k^{(0)}$. The algorithm stops when the convergence criterion is met. In the R package **ManlyMix**, we monitor the relative difference between Q-function values from two consecutive steps. If it is smaller than a user specified tolerance level, $1e-5$ by default, the algorithm stops. This is a speedier choice due to the fact that Q-function values are immediately available after numeric optimization of Equation 4. Such criterion is similar to monitoring the relative difference between log-likelihood values. Upon convergence, the Bayes decision rule assigns each observation to its cluster according to the maximized posterior probabilities from the last E-step. The estimated label of the i th observation is given by

$$\hat{Z}_i = \operatorname{argmax}_k \hat{\pi}_{ik}. \quad (5)$$

In **ManlyMix**, the function `Manly.EM()` runs the EM algorithm for a Manly mixture model and returns estimated model parameters, posterior probabilities, as well as a classification vector. This function is constructed in C for computational efficiency.

Pairwise overlap

Pairwise overlap, introduced by [Maitra and Melnykov \(2010\)](#), is a measure of the interaction between two mixture components. If we denote ω_{k_1, k_2} as the pairwise overlap of components k_1 and k_2 , it is defined as the sum of two misclassification probabilities

$$\omega_{k_1, k_2} = \omega_{k_1|k_2} + \omega_{k_2|k_1}, \quad (6)$$

where $\omega_{k_1|k_2}$ represents the probability that a random variable \mathbf{X} is mistakenly classified to group k_1 while it came from the component k_2 . For a Manly mixture, $\omega_{k_1|k_2}$ can be written as

$$\omega_{k_1|k_2} = \Pr \left[\frac{\phi(\mathcal{M}(\mathbf{X}; \lambda_{k_1}); \mu_{k_1}, \Sigma_{k_1}) \exp\{\lambda_{k_1}^T \mathbf{X}\}}{\phi(\mathcal{M}(\mathbf{X}; \lambda_{k_2}); \mu_{k_2}, \Sigma_{k_2}) \exp\{\lambda_{k_2}^T \mathbf{X}\}} > \frac{\tau_{k_2}}{\tau_{k_1}} \middle| \mathbf{X} \sim \phi(\mathcal{M}(\mathbf{x}; \lambda_{k_2}); \mu_{k_2}, \Sigma_{k_2}) \exp\{\lambda_{k_2}^T \mathbf{x}\} \right] \\ = \Pr \left[-\frac{1}{2} (\mathcal{M}(\mathcal{M}^{-1}(\mathbf{Y}; \lambda_{k_2}); \lambda_{k_1}) - \mu_{k_1})^T \Sigma_{k_1}^{-1} (\mathcal{M}(\mathcal{M}^{-1}(\mathbf{Y}; \lambda_{k_2}); \lambda_{k_1}) - \mu_{k_1}) + \lambda_{k_1}^T \mathcal{M}^{-1}(\mathbf{Y}; \lambda_{k_2}) \right. \\ \left. + \frac{1}{2} (\mathbf{Y} - \mu_{k_2})^T \Sigma_{k_2}^{-1} (\mathbf{Y} - \mu_{k_2}) - \lambda_{k_2}^T \mathcal{M}^{-1}(\mathbf{Y}; \lambda_{k_2}) > \log \left(\frac{\tau_{k_2} |\Sigma_{k_1}|^{1/2}}{\tau_{k_1} |\Sigma_{k_2}|^{1/2}} \right) \middle| \mathbf{Y} \sim \text{MVN}(\mu_{k_2}, \Sigma_{k_2}) \right]. \quad (7)$$

In **ManlyMix**, function `Manly.overlap()` estimates $\omega_{k_1|k_2}$ by sampling from corresponding distributions.

Variability assessment

The variability assessment of parameter estimates from Manly mixture model can be made by taking the inverse of the empirical observed information matrix $I_e(\hat{\Psi})$ ([McLachlan and Basford, 1988](#)) given by

$$I_e(\hat{\Psi}) = \sum_{i=1}^n \nabla q_i(\hat{\Psi}) \nabla q_i^T(\hat{\Psi}), \quad (8)$$

where $q_i(\Psi) = \sum_{k=1}^K \pi_{ik} [\log \tau_k + \log \phi(\mathcal{M}(x_i; \lambda_k); \mu_k, \Sigma_k) + \lambda_k x_i]$ and ∇ stands for the gradient operator. We take partial derivatives in the gradient vector $\nabla q_i(\Psi)$ and obtain

$$\begin{aligned} \frac{\partial q_i(\Psi)}{\partial \tau_k} &= \frac{\pi_{ik}}{\tau_k} - \frac{\pi_{iK}}{\tau_K}, & \frac{\partial q_i(\Psi)}{\partial \mu_k} &= \pi_{ik} \Sigma_k^{-1} (\mathcal{M}(x_i; \lambda_k) - \mu_k), \\ \frac{\partial q_i(\Psi)}{\text{vech}\{\Sigma_k\}} &= G^T \text{vec} \left\{ \frac{\pi_{ik}}{2} \Sigma_k^{-1} \left((\mathcal{M}(x_i; \lambda_k) - \mu_k)(\mathcal{M}(x_i; \lambda_k) - \mu_k)^T \Sigma_k^{-1} - I_p \right) \right\}, \\ \frac{\partial q_i(\Psi)}{\partial \lambda_k} &= -\pi_{ik} D_k \Sigma_k^{-1} (\mathcal{M}(x_i; \lambda_k) - \mu_k) + \pi_{ik} x_i, \end{aligned}$$

where I_p is the identity matrix of size p , $\text{vech}\{\cdot\}$ operator extracts the unique elements out of a symmetric $p \times p$ matrix and constructs a vector of length $p(p+1)/2$. G is a matrix with zero's and one's that enables the adoption of unique elements in a symmetric matrix (Melnykov, 2013). $\text{vec}\{\cdot\}$ is an operator which lines up all columns of a matrix one by one to form a vector. Finally,

$$D_k = \text{diag}\{(1 + (x_{i1}\lambda_{k1} - 1)e^{\lambda_{k1}x_{i1}})/\lambda_{k1}^2, \dots, (1 + (x_{ip}\lambda_{kp} - 1)e^{\lambda_{kp}x_{ip}})/\lambda_{kp}^2\}.$$

The estimated covariance matrix can be found as $I_e^{-1}(\hat{\Psi})$, i.e., by inverting the information matrix. Function `Manly.var()` in the package calculates the covariance matrix based on an estimated model provided by function `Manly.EM()`.

Forward and backward selection

In the Manly mixture model, there are $K \times p$ skewness parameters λ_{kj} corresponding to K components and p variables. Such a mixture is called a full Manly mixture model. Oftentimes, some coordinates are close to being normally distributed and the corresponding skewness parameters are unnecessary. Forward and backward selection procedures are adopted to eliminate such parameters and improve the model efficiency. These algorithms also prevent model-overfitting and conduct diagnostics with fitted skewness parameters. If underlying data groups are normally distributed, the selection procedures produce Gaussian mixture models.

The selection is based on the Bayesian information criterion (BIC) (Schwarz, 1978), which is the most commonly used criterion in finite mixture modeling (McLachlan and Peel, 2000). The smaller BIC is, the better fit provided by a mixture is. The forward selection procedure starts from the Gaussian mixture model and adds one λ_{kj} at a time until no improvement in BIC value can be obtained. The produced model is called Manly forward model (denoted as Manly F in this paper) with the details of the method outlined in Algorithm 1. The backward model selection algorithm given in Algorithm 2 works in the opposite direction. It starts with the full Manly mixture and drops one skewness parameter λ_{kj} at a time until no lower BIC can be reached. The obtained model is called

Data: X_1, \dots, X_n

Result: estimated model parameters by Manly forward model

Initialization: Gaussian mixture model

```

while the current model  $M_{\text{current}}$  has not reached the full Manly mixture model do
  1. find all zero skewness parameters in the current model  $M_{\text{current}}$ ,  $\lambda_1, \dots, \lambda_t$ ;
  2. construct new models  $M_{\text{new},1}, \dots, M_{\text{new},t}$  to compare with;
  3.  $M_{\text{new},j}$  sets the previous nonzero  $K \times p - t$  skewness parameters and  $\lambda_j$  to be
     non-zero;
  4. call function Manly.EM() to run the EM algorithm for each new model;
  5. initialize with the parameters of model  $M_{\text{current}}$  to speed the algorithm;
  if at least one new model has lower BIC than the original model  $M_{\text{current}}$  then
    find the smallest BIC among the new models;
    the corresponding new model  $M_{\text{new}}$  is selected and let  $M_{\text{current}} \leftarrow M_{\text{new}}$ .
  else
    break;
    the current model  $M_{\text{current}}$  is the final solution reached by Manly forward
    algorithm.
  end
end

```

Algorithm 1: Manly forward selection algorithm.

Data: X_1, \dots, X_n

Result: estimated model parameters by Manly backward model

Initialization: full Manly mixture model M_{full} with $K \times p$ non-zero skewness parameters

```

while the current model  $M_{current}$  has not reached Gaussian mixture model do
  1. find all non-zero skewness parameters in the current model  $M_{current}, \lambda_1, \dots, \lambda_s$ ;
  2. construct new models  $M_{new,1}, \dots, M_{new,s}$  to compare with;
  3.  $M_{new,j}$  sets the previous  $K \times p - s$  skewness parameters and  $\lambda_j$  to be zero;
  4. call function Manly.EM() to run the EM algorithm for each new model;
  5. initialize with the parameters of model  $M_{current}$  to speed the algorithm;
  if at least one new model has lower BIC than the original model  $M_{current}$  then
    find the smallest BIC among the new models; the corresponding new model
     $M_{new}$  is selected and let  $M_{current} \leftarrow M_{new}$ .
  else
    break;
    the current model  $M_{current}$  is the final solution reached by Manly backward
    algorithm.
  end
end

```

Algorithm 2: Manly backward selection algorithm.

the Manly backward model (Manly B). The selection algorithms are available in **ManlyMix** through setting `method = "forward"` or `method = "backward"` in the `Manly.select()` function.

Manly K-means clustering

Manly K-means clustering is constructed based on the classification EM (CEM) algorithm (Celeux and Govaert, 1992), which is a modification of the EM algorithm with an additional classification step. This step involves the Bayesian decision rule (i.e., $z_i^{(s)} = \operatorname{argmax}_k \pi_{ik}^{(s)}$) introduced immediately after the E-step.

It can be noticed that the traditional K-means algorithm is equivalent to the CEM algorithm based on the mixture model provided by

$$g(\mathbf{x}; \Psi) = \frac{1}{K} \sum_{k=1}^K \phi(\mathbf{x}; \mu_k, \sigma^2 \mathbf{I}).$$

The model underlying the traditional K-means imposes very restrictive assumptions of the homoscedasticity and spherical structure of components. We alleviate these assumptions by allowing each component to have the covariance matrix $\sigma_k^2 \mathbf{I}$ and applying Manly transformation to the data. These changes result in the model given by

$$g(\mathbf{x}; \Psi) = \frac{1}{K} \sum_{k=1}^K \phi(\mathcal{M}(\mathbf{x}; \lambda_k); \mu_k, \sigma_k^2 \mathbf{I}) \exp\{\lambda_k^T \mathbf{x}\}. \quad (9)$$

Following the same procedure as the Manly mixture EM algorithm, each λ_k can be obtained separately by straightforward numeric optimization of the function \tilde{Q}_k written as

$$\begin{aligned} \tilde{Q}_k(\lambda_k | \Psi^{(s-1)}) = & -\frac{pn_k^{(s)}}{2} \log \left\{ \sum_{i=1}^n \tilde{\zeta}_{ik}^{(s)} \left(n_k^{(s)} \mathcal{M}(\mathbf{x}_i; \lambda_k) - \sum_{j=1}^n \tilde{\zeta}_{jk}^{(s)} \mathcal{M}(\mathbf{x}_j; \lambda_k) \right) \right\}^T \\ & \times \left(n_k^{(s)} \mathcal{M}(\mathbf{x}_i; \lambda_k) - \sum_{j=1}^n \tilde{\zeta}_{jk}^{(s)} \mathcal{M}(\mathbf{x}_j; \lambda_k) \right) \left\{ \right\} + \lambda_k^T \sum_{i=1}^n \tilde{\zeta}_{ik}^{(s)} \mathbf{x}_i + \text{const}, \end{aligned}$$

where fuzzy classifications $\pi_{ik}^{(s)}$ are replaced by hard assignments in the form of indicators $\tilde{\zeta}_{ik}^{(s)} = I(z_i^{(s)} = k)$. If $z_i^{(s)} = k$ holds true, $\tilde{\zeta}_{ik}^{(s)}$ takes a value of 1; otherwise $\tilde{\zeta}_{ik}^{(s)}$ is equal to 0. The current size of the k th cluster is $n_k^{(s)} = \sum_{i=1}^n \tilde{\zeta}_{ik}^{(s)}$.

In this way, each step of the Manly K-means algorithm updates the partition and parameter

estimates. The partition update is given by

$$z_i^{(s)} = \underset{k}{\operatorname{argmin}} \left\{ \|\mathcal{M}(x_i; \lambda_k^{(s-1)}) - \mu_k^{(s-1)}\|^2 / (2(\sigma_k^2)^{(s-1)}) - (\lambda_k^{(s-1)})^T x_i + \frac{p}{2} \log(\sigma_k^2)^{(s-1)} \right\},$$

while the parameters are estimated through the following expressions:

$$\begin{aligned} \lambda_k^{(s)} &= \underset{\lambda_k}{\operatorname{argmax}} \tilde{Q}_k^{(s)}(\lambda_k), & \mu_k^{(s)} &= \sum_{i=1}^n \xi_{ik}^{(s)} \mathcal{M}(x_i; \lambda_k^{(s)}) / n_k^{(s)}, \quad \text{and} \\ (\sigma_k^2)^{(s)} &= \sum_{i=1}^n \xi_{ik}^{(s)} (\mathcal{M}(x_i; \lambda_k^{(s)}) - \mu_k^{(s)})^T (\mathcal{M}(x_i; \lambda_k^{(s)}) - \mu_k^{(s)}) / (pn_k^{(s)}). \end{aligned} \quad (10)$$

The Manly K -means algorithm is incorporated in the R package **ManlyMix** through the function `Manly.Kmeans()`. It can be used when the number of data points in each cluster is about the same and the transformed clusters are close to being spherical. It shows faster performance as the inversion of potentially large covariance matrices is not needed.

Package functionality and illustrative examples

All functions available in the package **ManlyMix** are listed with brief descriptions in Table 2. In this section, we demonstrate the utility of each function through a synthetic dataset and the analysis of two real-life datasets: *Iris* (Anderson, 1935; Fisher, 1936) and *AIS* (Cook and Weisberg, 1994).

Function	Description
<code>Manly.EM()</code>	Runs the EM algorithm for a Manly mixture model
<code>Manly.select()</code>	Runs forward and backward selection methods for a Manly mixture model
<code>Manly.Kmeans()</code>	Runs the Manly K -means clustering
<code>Manly.overlap()</code>	Estimates the overlap values for a Manly mixture
<code>Manly.sim()</code>	Simulates datasets from Manly mixture models
<code>Manly.var()</code>	Performs variability assessment of Manly mixture model parameter estimates and returns confidence intervals
<code>Manly.plot()</code>	Constructs a plot to display model-fitting and clustering
<code>ClassAgree()</code>	Calculates the confusion matrix and number of misclassifications
<code>Manly.model()</code>	Serves as a wrapper function for Manly mixture modeling

Table 2: Summary of functions implemented in **ManlyMix**.

Illustrative example 1

In this subsection, a Manly mixture is constructed with user-specified parameters. The overlap values of this mixture is estimated through function `Manly.overlap()`. Then function `Manly.sim()` simulates a dataset from the mixture along with a true membership vector.

Step a: Mixture specification

Now we demonstrate the procedure to construct a Manly mixture step by step. First, the user need to specify the number of components (assigned to K) and variables (assigned to p). In this case, we have a three-component bivariate mixture.

```
R> library(ManlyMix)
R> K <- 3; p <- 2; set.seed(123)
```

If the mixture probability density function of interest is written as

$$\begin{aligned} g(x) &= 0.25e^{0.2x_1+0.25x_2} \phi \left(\left(\frac{e^{0.2x_1}-1}{0.2}, \frac{e^{0.25x_2}-1}{0.25} \right); \begin{pmatrix} 4.5 \\ 7 \end{pmatrix}, \begin{pmatrix} 0.4 & 0 \\ 0 & 0.4 \end{pmatrix} \right) \\ &+ 0.3e^{0.5x_1+0.35x_2} \phi \left(\left(\frac{e^{0.5x_1}-1}{0.5}, \frac{e^{0.35x_2}-1}{0.35} \right); \begin{pmatrix} 4 \\ 8 \end{pmatrix}, \begin{pmatrix} 1 & -0.2 \\ -0.2 & 0.6 \end{pmatrix} \right) \\ &+ 0.45e^{0.3x_1+0.4x_2} \phi \left(\left(\frac{0.3e^{x_1}-1}{0.3}, \frac{e^{0.4x_2}-1}{0.4} \right); \begin{pmatrix} 5 \\ 5.5 \end{pmatrix}, \begin{pmatrix} 2 & -1 \\ -1 & 2 \end{pmatrix} \right), \end{aligned} \quad (11)$$

we construct the mixture by assigning the model parameter values to $\mathbf{1a}$ (matrix input of size $K \times p$), τ (vector input of length K), \mathbf{Mu} (matrix input of size $K \times p$) and \mathbf{S} (array input of dimensionality $p \times p \times K$), respectively.

```
R> tau <- c(0.25, 0.3, 0.45)
R> Mu <- matrix(c(4.5, 4, 5, 7, 8, 5.5), 3)
R> la <- matrix(c(0.2, 0.5, 0.3, 0.25, 0.35, 0.4), 3)
R> S <- array(NA, dim = c(p, p, K))
R> S[, , 1] <- matrix(c(0.4, 0, 0, 0.4), 2)
R> S[, , 2] <- matrix(c(1, -0.2, -0.2, 0.6), 2)
R> S[, , 3] <- matrix(c(2, -1, -1, 2), 2)
```

Step b: Overlap assessment

It is desirable to be capable of understanding the degree of interaction among mixing components to assess clustering complexity. Function `Manly.overlap()`, employing the measure of pairwise overlap, is implemented for this purpose. It has the following syntax:

```
Manly.overlap(tau, Mu, S, la, N = 1000)
```

with arguments \mathbf{la} , τ , \mathbf{Mu} , \mathbf{S} and N . Here, N represents the number of samples simulated from the given mixture for pairwise overlap estimation. The larger N is, the more precise the calculation is. By default, 1000 samples are employed. Four objects are returned by the function, including the misclassification probability matrix `$OmegaMap`, pairwise overlap `$OverlapMap`, average mixture overlap `$BarOmega`, and maximum mixture overlap `$MaxOmega`. Here, element `$OmegaMap[k2, k1]` corresponds to $\omega_{k_1|k_2}$ in Equation 7. In this case, for example, $\omega_{3|2} = 0.046$ means that a random variable coming from the second component has approximate probability of 0.046 to be misclassified to group 3. $\omega_{3|3} = 0.933$ represents the probability that a point belonging to group 3 is correctly assigned to this group. Each row of `$OmegaMap` sums up to 1. Then, pairwise overlaps ω_{k_1, k_2} given in Equation 6 are provided in the `$OverlapMap`. Among all pairwise overlaps ($\omega_{1,2}$, $\omega_{1,3}$ and $\omega_{2,3}$), $\omega_{2,3}$ yields the maximum value of 0.097 and produces `$MaxOmega`. The average of these three values, on the other hand, results in `$BarOmega` being 0.08066667.

```
R> A <- Manly.overlap(tau, Mu, S, la)
R> print(A)
$OmegaMap
      [,1] [,2] [,3]
[1,] 0.909 0.058 0.033
[2,] 0.038 0.916 0.046
[3,] 0.016 0.051 0.933

$OverlapMap
  Components Overlap
1      (1, 2)  0.096
2      (1, 3)  0.049
3      (2, 3)  0.097

$BarOmega
[1] 0.08066667

$MaxOmega
[1] 0.097
```

It can be seen that in the considered case, function `Manly.overlap()` calculates all characteristics based on the input of true model parameters. If parameters \mathbf{la} , τ , \mathbf{Mu} and \mathbf{S} are estimated, `Manly.overlap()` provides estimates of misclassification probabilities and overlap values. As for high-dimensional data, we can not readily visually assess the interaction between data groups, such output helps approximate the proximity of clusters and discover properties associated with them.

Step c: Data generation

Function `Manly.sim()` simulates Manly mixture datasets based on user-specified model parameters. It employs the built-in R function `rmultinom()` for assigning data points to K mixture components according to the mixing proportion τ_k 's. Then the function simulates normally distributed data points by function `rnorm()`. The covariance structures Σ_k are applied to the data points before back-transforming them to Manly distributed components.

The `Manly.sim()` command has the following syntax:

```
Manly.sim(n, la, tau, Mu, S).
```

The user can input n as the desired sample size. Here, a dataset of 30 observations is simulated from Equation 11 and data matrix X as well as its true membership vector $\$id$ are returned.

```
R> n <- 30
R> B <- Manly.sim(n, la, tau, Mu, S)
R> print(B)
$X
      [,1]      [,2]
[1,] 3.259485 3.882271
[2,] 3.247362 4.269247
```

Part of the output is intentionally omitted.

```
[30,] 3.310186 2.974554

$id
[1] 1 1 1 1 1 1 1 2 2 2 2 2 2 2 2 2 2 2 3 3 3 3 3 3 3 3 3 3
```

Illustrative example 2: *Iris* dataset

Iris dataset (Anderson, 1935; Fisher, 1936) has 150 observations and 4 variables that represent *sepal length*, *sepal width*, *petal length*, and *petal width*. Three clusters, *Setosa*, *Versicolor*, and *Virginica*, are of equal size, consisting of 50 observations each. The function `Manly.EM()` fits a Manly mixture to the *Iris* dataset and 95% confidence intervals of the model MLE are provided by `Manly.var()`. The Manly F and Manly B models are obtained by `Manly.select()`. The Manly K-means algorithm clusters the dataset through `Manly.Kmeans()`.

Step a: Data preparation

`Manly.EM()` requires input of a matrix object X , where rows of X represent p -variate observations. If X is univariate data with vector input, it will be automatically transformed into a matrix of just one column. Thus, X has the dimensionality $n \times p$. In this case, we transform the *Iris* dataset into a matrix of dimensionality 150×4 and assign it to X .

```
R> library(ManlyMix)
R> K <- 3; p <- 4
R> X <- as.matrix(iris[, -5])
```

Step b: Initialization of the EM algorithm

Good initialization strategy of the EM algorithm is important to improve chances of finding a correct result. There are two ways for the user to initialize the `Manly.EM()` function. One is by means of providing the initial partition of the data id (vector input of length n) and skewness parameters la . Here, it needs to be noticed that the specification of la matrix serves as an indicator of whether the transformation is applied to a specific variable and component or not. For example, for the *Iris* dataset, assumes that all variables in all components enjoy normality except for the first variable in the first component, la needs to be set as $la \leftarrow \text{matrix}(c(0.1, \text{rep}(0, 11)), 3, 4)$, with 0.1 (that can be any non-zero value) serving as the starting point in Nelder-Mead optimization. If no la is provided, the skewness parameters are all set equal to 0 and a Gaussian mixture model will be fitted. The other way of starting the algorithm is to enter initial model parameters, including la , τ , μ , and S . The algorithm employs these parameters to compute the posterior probabilities in the first E-step.

Here, we adopt the first strategy. The initial partition of the *Iris* data is obtained by running the traditional K-means algorithm and specifying la is a matrix of size 3×4 , with all elements set to a non-zero value of 0.1.

```
R> set.seed(123)
R> id.km <- kmeans(X, K)$cluster
R> la <- matrix(0.1, K, p)
```

Step c: EM algorithm for Manly mixture modeling

`Manly.EM()` runs the EM algorithm for modeling based on Manly mixtures given in Equation 1. The command has the following syntax:

```
Manly.EM(X, id = NULL, la = NULL, tau = NULL, Mu = NULL, S = NULL,
tol = 1e-5, max.iter = 1000).
```



```
R> G <- Manly.EM(X, id = id.km)
R> colnames(G$la) <- colnames(X)
R> print(G$la)
      Sepal.Length Sepal.Width Petal.Length Petal.Width
[1,]              0              0              0              0
[2,]              0              0              0              0
[3,]              0              0              0              0

R> M <- Manly.EM(X, id.km, la)
R> colnames(M$la) <- colnames(X)
R> print(M$la)
      Sepal.Length Sepal.Width Petal.Length Petal.Width
[1,]   -0.1158602   0.05907443   -0.2382086   -4.033529
[2,]   -0.1254022  -0.65079974   -0.3848938    0.479587
[3,]   -0.1282339   0.64271380    0.3343054   -1.134275

R> print(M$id)
      [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
      [38] 1 1 1 1 1 1 1 1 1 1 1 1 3 3 3 3 3 3 3 3 3
      [75] 3 3 3 3 3 3 3 3 2 3 3 3 3 3 3 3 3 3 3 3 3
      [112] 2 2 2 2 2 2 2 2 3 2 2 2 3 2 2 3 3 2 2 2 2 3 2
      [149] 2 2
```

The characteristics of the fitted model are demonstrated in terms of the model log-likelihood `$ll` and BIC `$bic`. The number of iterations run by the EM algorithm until convergence is recorded through `$iter`. In this example, the EM algorithm reaches convergence after 13 iterations and the model BIC is 618.46. Finally, a dummy indicator `$flag` reports the validity of the fitted model, where 0 represents the successful convergence of the EM algorithm and 1 stands for the failure of convergence. A warning message is given if `$flag` is equal to 1. It may happen when one cluster disappears or shrinks so that some parameter estimates are NA's. Such issue is related to spurious solutions (McLachlan and Peel, 2000) where one or more components model a local pattern in data rather than a systematic one.

X represents the data matrix and `model` is the object of class "ManlyMix". `conf.CI` is user-specified confidence level, which needs to take a value between 0 and 1. Here by setting `model = M`, we take the MLE of the fitted Manly mixture obtained from step c and evaluate its variability. The number of unique model parameters is $K - 1 + 2K \times p + K \times p(p + 1)/2 = 56$ for the *Iris* dataset. `conf.CI =`

0.95 calculates 95% confidence intervals for these 56 parameters. Thus `Manly.var()` function returns a 56×56 covariance matrix (assigned to `V`) and 56 confidence intervals (assigned to `CI`).

```
R> result <- Manly.var(X, model = M, conf.CI = 0.95)
```

In the code output of 95% confidence intervals, the first column represents the point estimates of the 56 model parameters, while the second and third columns stand for the lower and upper bounds of confidence intervals, respectively.

```
R> print(result$CI)
      Estimates      Lower      Upper
[1,] 0.33333333 0.257887628 0.408779039
[2,] 0.264119102 0.175676489 0.352561716
[3,] 3.794594378 -8.303528084 15.892716840
```

Part of the output is intentionally omitted.

```
[54,] 0.642713799 -0.407079526 1.692507125
[55,] 0.334305435 -0.128841410 0.797452280
[56,] -1.134275340 -2.079185136 -0.189365544
```

Step e: Forward and backward selection algorithms

Step e targets detecting the normally distributed variables in *Iris*. `Manly.select()` provides the selection algorithm for eliminating unnecessary skewness parameters in M\$1a. These skewness parameters are fixed to be equal to zero and the log-likelihood is maximized based on the rest of parameters. The use of the function is shown below:

```
Manly.select(X, model, method, tol = 1e-5, max.iter = 1000, silent = FALSE).
```

The argument `model` is the initial model to start the selection procedure with. `method` is set to either "forward" or "backward" for the implementation of Algorithm 1 or Algorithm 2, respectively. The selection criterion for each step is based on `$bic` values obtained from all candidate models that are of class "ManlyMix". `silent` is an argument that controls the code output. By default, `silent` provides the steps of selection and BIC values for all candidate models. Thus, the user can monitor the selection procedures. The output can be turned off by setting `silent = TRUE`. We first discuss the implementation of the forward selection on the *Iris* dataset. The algorithm is initialized by the Gaussian mixture model G obtained in step c.

```
R> MF <- Manly.select(X, model = G, method = "forward")
step 1 :
  current BIC = 580.8389
  alternative BICs = 585.6791 585.0607 582.1893 585.7369 585.2193 583.7374
                    585.7081 583.963 579.8978 573.4626 585.8161 585.8407
step 2 :
  current BIC = 573.4626
  alternative BICs = 578.3191 577.6844 574.813 578.3719 577.843 576.3611
                    578.3282 576.5866 572.5215 578.4397 578.4643
step 3 :
  current BIC = 572.5215
  alternative BICs = 577.378 576.7713 575.8067 577.4308 576.8833 575.3526
                    577.3871 575.6213 577.4799 577.3221
```

The forward selection takes three steps for the algorithm to find the best model (assigned to MF). In step 3, there is no alternative BIC value that is smaller than the current model BIC, so the forward selection algorithm stops searching over non-zero λ_{kj} 's. Compared to the Gaussian mixture fit, Manly F model improves by 8 in BIC value.

On the contrary, the backward selection starts with the full Manly mixture M and drops one skewness parameter at a time.

```
R> MB <- Manly.select(X, model = M, method = "backward")
step 1 :
  current BIC = 618.4553
  alternative BICs = 613.5161 612.7184 613.8658 613.448 614.3828 616.6445
                    613.5442 616.3626 617.0157 625.7431 613.1927 610.7879
step 2 :
  current BIC = 610.7879
  alternative BICs = 605.9157 605.9075 607.3512 605.8475 606.3851 607.8112
                    605.9437 607.0513 609.9457 618.1426 605.7915
```

Part of the output is intentionally omitted.

```
step 10 :
  current BIC = 575.3526
  alternative BICs = 572.5215 576.3611 582.729
step 11 :
  current BIC = 572.5215
  alternative BICs = 573.4626 579.8978
```

After 11 steps, the backward selection produces the Manly B model, which enjoys the same BIC value as the Manly F model.

Step f: Diagnostics

The skewness parameters of the Manly F and Manly B models are investigated in the following example. It is observed that the forward selection adopts only two λ_{kj} 's in the model. They correspond to the *petal width* variable of the first species and the *petal length* variable of the third one. For all other components and variables, the data appear to be nearly normally distributed. The same two skewness parameters are found by the backward selection. It is worth mentioning, however, that Manly F and Manly B models can produce different results.

```
R> colnames(MF$la) <- colnames(X)
R> print(MF$la)
  Sepal.Length Sepal.Width Petal.Length Petal.Width
[1,]          0          0    0.0000000    -4.04
[2,]          0          0    0.0000000     0.00
[3,]          0          0    0.5615625     0.00
R> colnames(MB$la) <- colnames(X)
R> print(MB$la)
  Sepal.Length Sepal.Width Petal.Length Petal.Width
[1,]          0          0    0.0000000   -4.034815
[2,]          0          0    0.0000000     0.000000
[3,]          0          0    0.5619671     0.000000
```

Step g: Manly K-means algorithm

The Manly K-means algorithm written in Equation 9 is implemented in function `Manly.Kmeans()`, which has the following syntax:

```
Manly.Kmeans(X, id = NULL, la = NULL, Mu = NULL, S = NULL,
  initial = "k-means", K = NULL, nstart = 100, method = "ward.D",
  tol = 1e-5, max.iter = 1000).
```

`Manly.Kmeans()` has most of the arguments and returned values the same as those of the function `Manly.EM()`. As the Manly K-means algorithm assumes that all clusters are of the same size, the mixing proportions τ are not needed in this function. S is a vector of length K that represents variance within each cluster, as the transformed data groups are assumed to be spherical. The parameters returned by the function $\$la$, $\$Mu$, and $\$S$ correspond to λ_k , μ_k and σ_k^2 given in Equation 10. The log-likelihood and BIC values are not provided since the parameter estimates are not MLE's. `Manly.Kmeans()` has several initialization choices: (1) by providing `id` and `la`; (2) by providing `la`, `Mu`, and `S`; (3) by specifying the number of clusters K and letting `initial = "k-means"`; It takes the default traditional K-means clustering result and passes it into the CEM algorithm; (4) by specifying the number of clusters K and letting `initial = "hierarchical"`; It adopts the hierarchical clustering solution as the initial dataset partition. `nstart` is responsible for controlling the number of random starts tried in initialization choice (3) with a default value equal to 100. `method` sets the linkage method in initialization choice (4) with a default of `method = "ward.D"`, which represents the Ward's linkage (Ward, 1963). Here, the initialization choice of `Manly.Kmeans()` is (1), which is the same as that of `Manly.EM()`.

```
R> MK <- Manly.Kmeans(X, id.km, la)
R> colnames(MK$la) <- colnames(X)
R> print(MK$la)
  Sepal.Length Sepal.Width Petal.Length Petal.Width
[1,] -0.37975529 -0.5815382  -0.81530022  -2.5572830
[2,] -0.27067058 -0.4103692  -0.31001602  -0.5367999
[3,] -0.02896526  0.1138177  -0.05694487   0.2617650

R> print(MK$S)
[1] 0.002717844 0.006156015 0.160435910
```

Illustrative example 3: AIS dataset

In this subsection, dataset *AIS* (Cook and Weisberg, 1994) is studied for illustrative purposes. The Australian Institute of Sports (*AIS*) dataset was first introduced by Cook and Weisberg (1994). It contains information collected from 202 athletes, among which 100 are females and 102 are males. There are 13 variables, including the gender, sport kind and 11 numeric measurements of the athletes. We adopt the same variables and analysis as Lee and McLachlan (2013). The goal of the analysis is to cluster the athletes into two groups: males and females by constructing models based on three measurements: the body mass index (*BMI*), lean body mass (*LBM*), and the percentage of body fat (*Bfat*). Function `ClassAgree()` compares the estimated and true partitions. Function `Manly.plot()` is introduced for visual analysis of Manly mixture fitted results.

Step a: model fit

The *AIS* dataset is analyzed by six mixture models: the traditional *K*-means (`kmeans()`), Manly *K*-means (`Manly.Kmeans()`), Gaussian mixture model (`Manly.EM()`), Manly mixture model (`Manly.EM()`), Manly forward model and Manly backward model (both available through `Manly.select()`).

```
R> library(ManlyMix)
R> data("ais"); set.seed(123)
R> X <- as.matrix(ais[,c(8,10,11)]); id <- as.numeric(ais[,1])
R> n <- dim(X)[1]; p <- dim(X)[2]; K <- max(id)
R> Kmeans <- kmeans(X, K); id.km <- Kmeans$cluster
```

By running the following code, we not only obtain the fitted models, but also test the package from different aspects. The number of parameters in the models are 7 (*K*-means), 19 (Gaussian), 25 (Manly), 23 (Manly F), 22 (Manly B) and 14 (Manly *K*-means). The computing times are 0.001, 0.004, 0.083, 0.8, 1.143, 0.024, respectively. These results are rather efficient compared to those from other packages (see Appendix).

```
R> MK <- Manly.Kmeans(X, id = id.km, la = matrix(0.1, K, p))
R> G <- Manly.EM(X, id = id.km, la = matrix(0, K, p))
R> M <- Manly.EM(X, id = id.km, la = matrix(0.1, K, p))
R> MF <- Manly.select(X, G, method = "forward", silent = TRUE)
R> MB <- Manly.select(X, M, method = "backward", silent = TRUE)
```

Now we consider the fitted model parameters to perform a comprehensive analysis and diagnostics of the *AIS* dataset. From the following output, it is observed that the Manly F model drops two skewness parameters from the full Manly mixture model while Manly B drops three. This yields the conclusion that the *Bfat* variable in the first group and *LBM* variable in the second one are close to be normal. Through the one-to-one correspondence between skewness parameters and dataset variables, **ManlyMix** is proved to be particularly useful for model variable diagnostics.

```
R> colnames(MF$la) <- colnames(X)
R> MF$la
      BMI      Bfat      LBM
[1,] -0.08671894  0.0000000 0.01002851
[2,] -0.12882354 -0.1902031 0.00000000

R> colnames(MB$la) <- colnames(X)
R> MB$la
      BMI      Bfat LBM
[1,] -0.09362427  0.0000000  0
[2,] -0.12720459 -0.1933216  0
```

BIC values for the four models are 3595.35 (Gaussian), 3543.00 (Manly), 3538.42 (Manly F) and 3533.63 (Manly B). It shows considerable improvement in terms of BIC from Manly mixture models. They provide better fits for the data, among which Manly backward is the best model selected according to BIC.

Step b: classification table

Classification results from the six models are compared using function `ClassAgree()` in step b. Function `ClassAgree()` adopts input of both the estimated and true id vectors with the following syntax:

```
ClassAgree(est.id, trueid).
```

`ClassAgree()` permutes the partition labels to achieve the lowest number of misclassifications. Then, based on the switched labels, it returns the confusion matrix and number of misclassifications. In the analysis of the *AIS* dataset, the following output is produced by `ClassAgree()`.

```
R> ClassAgree(id.km, id)
$ClassificationTable
  est.id
trueid 1 2
      1 98 2
      2 12 90
```

```
$MisclassificationNum
[1] 14
```

```
R> ClassAgree(MK$id, id)
$ClassificationTable
  est.id
trueid 1 2
      1 95 5
      2 7 95
```

```
$MisclassificationNum
[1] 12
```

```
R> ClassAgree(G$id, id)
$ClassificationTable
  est.id
trueid 1 2
      1 100 0
      2 8 94
```

```
$MisclassificationNum
[1] 8
```

```
R> ClassAgree(M$id, id)
$ClassificationTable
  est.id
trueid 1 2
      1 98 2
      2 2 100
```

```
$MisclassificationNum
[1] 4
```

```
R> ClassAgree(MF$id, id)
$ClassificationTable
  est.id
trueid 1 2
      1 99 1
      2 3 99
```

```
$MisclassificationNum
[1] 4
```

```
R> ClassAgree(MB$id, id)
$ClassificationTable
  est.id
trueid 1 2
      1 99 1
      2 4 98
```

```
$MisclassificationNum
[1] 5
```

Rows and columns represent the true and estimated partitions, respectively. The diagonal and off-diagonal elements in the table correspond to correct and incorrect classifications, respectively. The lowest number of misclassifications (4 misclassifications) is obtained by the Manly mixture and Manly forward models. One worth-mentioning fact is that these two models enjoy the clustering solution as

good as the unrestricted skew- t mixture, which is reported to be the best model by [Lee and McLachlan \(2013\)](#). The Manly backward model comes second with 5 misclassifications. The remaining three models, traditional K -means, Manly K -means and Gaussian mixture model show worse performance.

Step c: visualization tool

In order to investigate the behavior of each model, contour plots with classified data points need to be analyzed. `Manly.plot()` allows conducting the visual analysis of a dataset fitted by Manly mixture model. The command has the following syntax:

```
Manly.plot(X, var1 = NULL, var2 = NULL, model = NULL, x.slice = 100,
y.slice = 100, x.mar = 1, y.mar = 1, col = "lightgrey", ...).
```

If both `var1` and `var2` are provided, they represent variables on the X-axis and Y-axis of a contour plot, respectively. Argument `model` is the object of class "ManlyMix". The parameters of `model` object are used to calculate the density and draw contour lines. The estimated membership vector `model$id` is reflected through different colors. `x.slice` and `y.slice` options control the number of grid points for which a density is calculated. The larger these two values are, the more grid values are considered. Thus, the contour lines look smoother. `x.mar` and `y.mar` specify plot margins. The parameter `col` specifies the color of contour lines with the default color being light grey. Other variables in the built-in R function `contour()` can also be used as specified. On the other hand, if only `var1` is provided, a density plot of this variable is constructed. `x.slice` and `x.mar` have the same functionality as those in the contour plot. The parameter `col` stands for density line color with the default being light grey. ... allows other arguments from the built-in R function `hist()` to be passed.

In this case, we conduct the same analysis as that in [Lee and McLachlan \(2013\)](#) and adopt the two variables *LBM* and *Bfat* for constructing contour plots. The margins of the plots are set to be 3 on the X-axis and 13 on the Y-axis. The light grey contour lines have width equal to 3.2. Labels and axes are suppressed. The function is first applied to the four fitted models Gaussian mixture, Manly mixture, Manly F and Manly B in step a.

```
R> Manly.plot(X, var1 = 3, var2 = 2, model = G, x.mar = 3, y.mar = 13,
+ xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "", nlevels = 10,
+ drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
R> Manly.plot(X, var1 = 3, var2 = 2, model = M, x.mar = 3, y.mar = 13,
+ xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "", nlevels = 10,
+ drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
R> Manly.plot(X, var1 = 3, var2 = 2, model = MF, x.mar = 3, y.mar = 13,
+ xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "", nlevels = 10,
+ drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
R> Manly.plot(X, var1 = 3, var2 = 2, model = MB, x.mar = 3, y.mar = 13,
+ xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "", nlevels = 10,
+ drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
```

Function `Manly.plot()` enjoys sufficient flexibility to adopt other parsimonious models. Parameters obtained by traditional K -means and Manly K -means can be adjusted according to the object of class "ManlyMix" so that `$id`, `$tau`, `$Mu`, `$la`, `$S` are extracted in their correct forms.

```
R> Kmeans$id <- id.km; Kmeans$tau <- MK$tau <- rep(1/K, K)
R> Kmeans$Mu <- Kmeans$centers; Kmeans$la <- matrix(0, K, p)
R> Kmeans$S <- array(0, dim = c(p,p,K))
R> for(k in 1:K){diag(Kmeans$S[,k]) <- Kmeans$tot.withinss / n / p}
R> s2 <- MK$S; MK$S <- array(0, dim = c(p,p,K))
R> for(k in 1:K){diag(MK$S[,k]) <- s2[k]}
R> Manly.plot(X, var1 = 3, var2 = 2, model = Kmeans, x.mar = 3, y.mar = 13,
+ xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "", nlevels = 4,
+ drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
R> Manly.plot(X, var1 = 3, var2 = 2, model = MK, x.mar = 3, y.mar = 13,
+ xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "", nlevels = 10,
+ drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
```

Figure 1 combines all output plots from `Manly.plot()`.

It can be observed that the two components have a slight overlap, so the clustering problem is not over-complicated. However, the red cluster is highly skewed and has a heavy tail. This imposes difficulties for the traditional K -means, Manly K -means, and Gaussian mixture model. Manly mixture model shows great flexibility and captures the skewness pattern in both components. Manly forward drops the skewness parameters associated with variable *Bfat* in the female cluster (black component) and *LBM* in the male group (red component). Manly backward drops both skewness parameters that

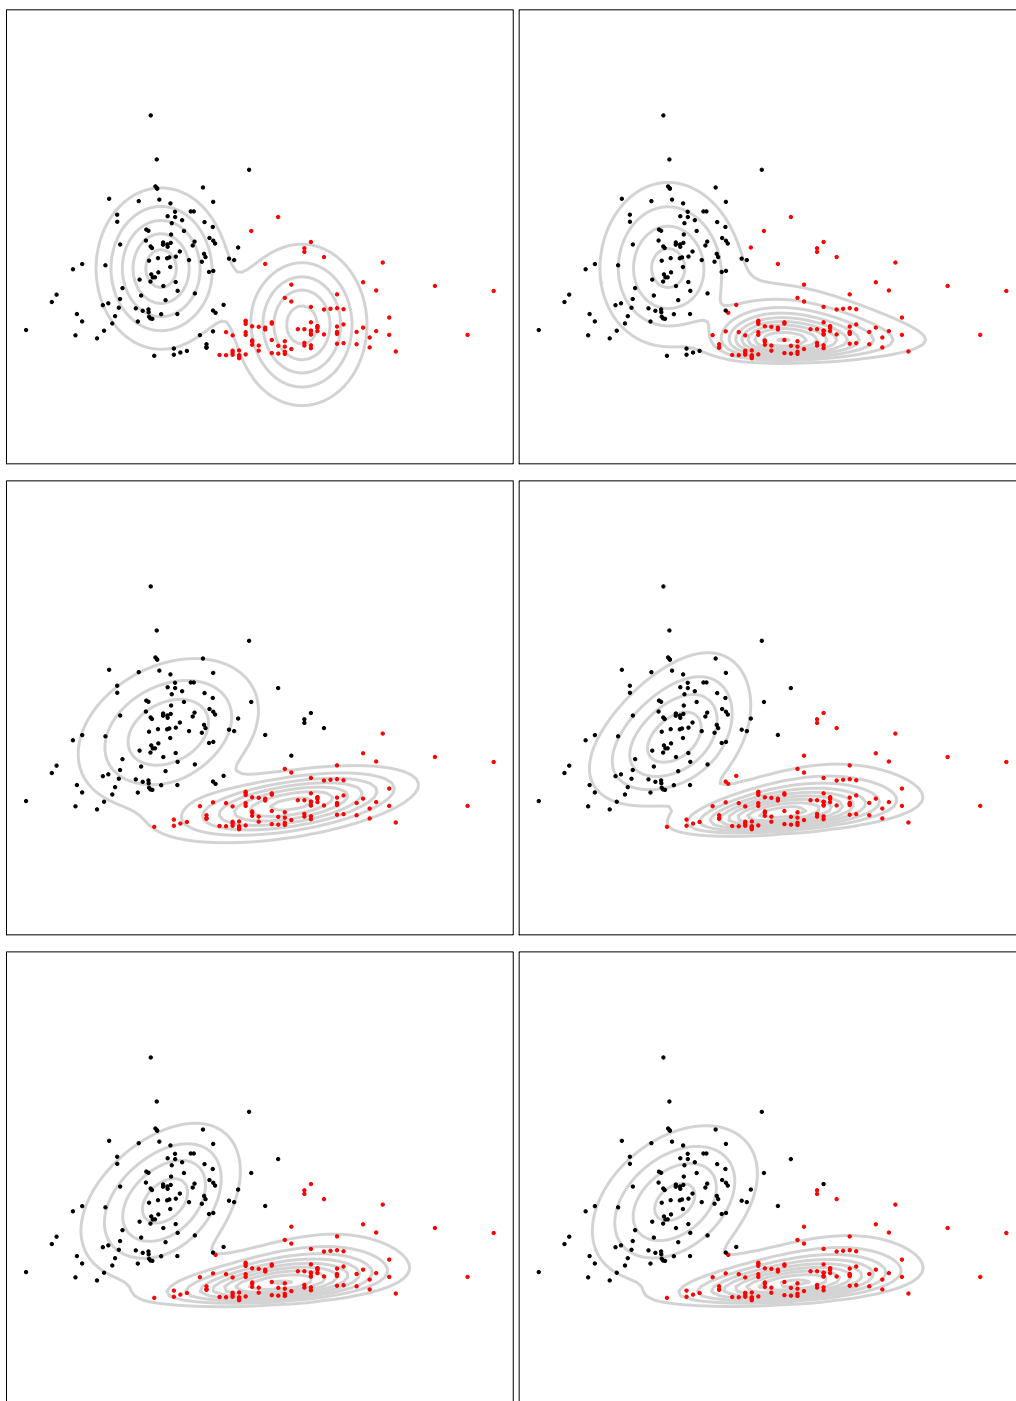


Figure 1: AIS dataset: fitted contour plots from function `Manly.plot()` based on the two variables *LBM* (X-axis) and *Bfat* (Y-axis). The model locations are: *K*-means (first row left), Manly *K*-means (first row right), Gaussian mixture (second row left), Manly mixture (second row right), Manly forward (third row left) and Manly backward (third row right).

correspond to the female cluster and uses a black ellipsoid. It also drops the *LBM* variable in the male cluster (red component). The above results reveal the applicability and effectiveness of function `Manly.plot()` on real-life datasets.

Alternative coding d: wrapper function

Wrapper function `Manly.model()` enables practitioners to run analysis in a simple and convenient way. The function has the following syntax:

```
Manly.model(X, K = 1:5, Gaussian = FALSE, initial = "k-means",
  nstart = 100, method = "ward.D", short.iter = 5, select = "none",
  silent = TRUE, plot = FALSE, var1 = NULL, var2 = NULL,
  VarAssess = FALSE, conf.CI = NULL, overlap = FALSE, N = 1000, tol = 1e-5,
  max.iter = 1000, ...).
```

Argument *K* is an integer vector providing the numbers of clusters to be tested for the data. The default setting tests 1, 2, 3, 4, or 5 clusters. It calls the `Manly.EM()` function to fit all five models. The one with the lowest BIC value is chosen to be the best model. Gaussian option specifies whether skewness parameters are adopted or not. If `TRUE`, Gaussian mixtures are fitted. With the default value being `FALSE`, it runs full Manly mixture models. *initial* specifies the initialization strategy used. It has three input options: (1) *initial* = "k-means" is the default initialization strategy, which passes the traditional *K*-means result into the EM algorithm as the initial partition. *nstart* is passed into the built-in R function `kmeans` for specifying the number of random starts (the default *nstart* = 100); (2) if *initial* = "hierarchical", the hierarchical clustering initialization is used. The linkage method is passed by *method* argument into R function `hclust`. The default is Ward's linkage; (3) if *initial* = "emEM", the emEM (Biernacki et al., 2003) initialization is run. Short runs of EM are conducted based on random starts and the one that corresponds to the highest log-likelihood is picked for running until convergence. *nstart* controls the number of random starts. The number of iterations for the short EM is specified by *short.iter* with a default value set to 5 iterations.

select argument has three input values: "none", "forward" and "backward". If *select* = "none", then the object returned by function `Manly.EM()` is adopted directly. If *select* = "forward", the Gaussian option is automatically adjusted to `Gaussian = TRUE`. It calls function `Manly.select(..., method = "forward")` to improve the original Gaussian fit. On the other hand, if *select* = "backward", Gaussian option is automatically set to `Gaussian = FALSE`. The full Manly mixture is followed by the backward selection `Manly.select(..., method = "backward")`. *silent* argument controls the output in function `Manly.select()`. The default setting suppresses the output. *plot* determines whether `Manly.plot()` function is called or not. If *plot* = `TRUE`, then `Manly.plot()` runs and arguments *var1* and *var2* allow user to specify which variable(s) to plot. Argument *VarAssess* provides the option of using `Manly.var()` for variability assessment. Notice here that it only provides assessment for a full Manly mixture model. *conf.CI* specifies the confidence level of the confidence intervals returned. *overlap* option, if specified to be `TRUE`, adopts `Manly.overlap()` function and estimates pairwise overlap values for the returned model. *N* is the number of Monte Carlo simulations run in `Manly.overlap()`.

Three objects are returned by function `Manly.model()`: `$model`, `$VarAssess`, and `$Overlap`. `$model` is the final model of class "ManlyMix" by `Manly.EM()` or `Manly.select()`. `$VarAssess` returns the variance-covariance matrix and confidence intervals by `Manly.var()` function. `$Overlap` returns the object by `Manly.overlap()`.

For *AIS* dataset, suppose the user wants to obtain the Manly F or Manly B model and take a look at their contour plots. A compact version of the code is given by:

```
R> MF <- Manly.model(X, K = 2, initial = "k-means", nstart = 100,
+ select = "forward", plot = TRUE, var1 = 3, var2 = 2, x.mar = 3,
+ y.mar = 13, xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "",
+ nlevels = 4, drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
R> MB <- Manly.model(X, K = 2, initial = "k-means", nstart = 100,
+ select = "backward", plot = TRUE, var1 = 3, var2 = 2, x.mar = 3,
+ y.mar = 13, xaxs="i", yaxs="i", xaxt="n", yaxt="n", xlab="", ylab = "",
+ nlevels = 4, drawlabels = FALSE, lwd = 3.2, col = "lightgrey", pch = 19)
```

Here, `MF$model` and `MB$model` obtained are the same as those in step a. The contour plots are generated automatically and can be found in Figure 1.

Alternative coding e: initialization with model parameters

Functions `Manly.EM()` can take initial model parameters as initialization of the algorithm. It is especially useful for the emEM initialization. The practitioner can construct a large number of short EM runs, select the one with the highest log-likelihood and obtain its estimated parameters. Then, the EM algorithm initialized by these parameters is run until convergence. Here is a small example on the

AIS dataset. 100 short EM algorithms run for 5 iterations each. As we can see, the obtained object M is the same as that from step a.

```
R> ll <- -Inf; init <- NULL
R> nstart <- 100; iter <- 0
R> repeat{
+   id.km <- kmeans(X, centers = K, iter.max = 1)$cluster
+   temp <- Manly.EM(X, id = id.km, la = matrix(0.1, K, p), max.iter = 5)
+   if(temp$ll > ll){
+     ll <- temp$ll
+     init <- temp
+   }
+   iter <- iter + 1
+   if(iter == nstart){break}
+ }
R> M <- Manly.EM(X, tau = init$tau, Mu = init$Mu, S = init$S, la = init$la)
```

Illustrative example 4: acidity dataset

Since one reviewer is interested in seeing a showcase of a univariate Manly mixture, we illustrate its utility on the *acidity* dataset (Crawford, 1994). It provides the acidity measure of 155 lakes in the Northeastern United States. There are two clusters, but the true partition is unknown.

Step a: model fit

We run the following models: the traditional K-means (`kmeans()`), Manly K-means (`Manly.Kmeans()`), Gaussian mixture model (`Manly.EM()`), Manly mixture model (`Manly.EM()`), Manly forward model, and Manly backward model (both available through `Manly.select()`).

```
R> library(ManlyMix)
R> set.seed(123); data("acidity")
R> K <- 2; p <- 1; X <- acidity
R> Kmeans <- kmeans(X, K); id.km <- Kmeans$cluster
R> MK <- Manly.Kmeans(X, id = id.km, la = matrix(0.1, K, p))
R> G <- Manly.EM(X, id = id.km, la = matrix(0, K, p))
R> M <- Manly.EM(X, id = id.km, la = matrix(0.1, K, p))
R> MF <- Manly.select(X, G, method = "forward", silent = TRUE)
R> MB <- Manly.select(X, M, method = "backward", silent = TRUE)
```

The model BIC values for Gaussian, Manly, Manly F and Manly B are 394.51, 389.84, 389.84 and 389.84, respectively. There is an indication of skewness as both the Manly F and Manly B models fail to drop any skewness parameters. The Manly models improve by 5 in the BIC value.

Step b: visualization tool

To visually assess the fit provided by all models, we use the command `Manly.plot()` with univariate input. The fitted density plots associated with histogram of the data are provided in Figure 2.

```
R> Kmeans$id <- id.km; Kmeans$tau <- MK$tau <- rep(1/K, K)
R> Kmeans$Mu <- Kmeans$centers; Kmeans$la <- matrix(0, K, p)
R> Kmeans$S <- array(0, dim = c(p,p,K))
R> for(k in 1:K){Kmeans$S[,k] <- Kmeans$tot.withinss / n / p}
R> s2 <- MK$S; MK$S <- array(0, dim = c(p,p,K))
R> for(k in 1:K){MK$S[,k] <- s2[k]}
R> Manly.plot(X= acidity, model = Kmeans, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
R> Manly.plot(X= acidity, model = MK, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
R> Manly.plot(X= acidity, model = G, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
R> Manly.plot(X= acidity, model = M, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
R> Manly.plot(X= acidity, model = MF, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
R> Manly.plot(X= acidity, model = MB, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
```

Manly models provide the most reasonable fit of the data. The first component is slightly skewed to the right and only the Manly models pick up the high density at its peak. The second component is

slightly skewed to the left. The density fits provided by *K*-means and Manly *K*-means are insufficient due to the assumption of equal size components.

Alternative coding c: wrapper function

The wrapper function `Manly.model()` is capable of combining steps a and b in one command. The following code directly yields the Manly F or Manly B model:

```
R> MF <- Manly.model(X, K = 2, Gaussian = TRUE, initial = "k-means",
+ nstart = 100, select = "forward", plot = TRUE, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
R> MB <- Manly.model(X, K = 2, Gaussian = FALSE, initial = "k-means",
+ nstart = 100, select = "backward", plot = TRUE, var1 = 1, main = "", ylim = c(0, 0.75),
+ xlab = "", xaxt = "n", ylab = "", yaxt = "n", x.slice = 200, col = "red")
```

The Manly F and Manly B density plots given in Figure 2 are generated automatically.

Demo examples

For users who need further information about the package, we have constructed 16 demo examples listed in Table 4 that provide a comprehensive demonstration of **ManlyMix** capabilities. Among the examples, 11 of them are designed to demonstrate the capability and utility of each function and 5 of them run comprehensive analysis of classification datasets. Each demo can be accessed by its name and the users can reproduce themselves.

As an illustration of how these demos can be employed, the code of the first example can be approached through running the following code in R.

```
library(ManlyMix)
demo(EMalgorithm1)
```

Function	Demo example(s)
Manly.EM()	demo(EMalgorithm1), demo(EMalgorithm2)
Manly.select()	demo(ForwardSelection), demo(BackwardSelection)
Manly.Kmeans()	demo(ManlyKmeans1), demo(ManlyKmeans2)
Manly.overlap()	demo(Overlap)
Manly.sim()	demo(DataSimulation)
Manly.var()	demo(VarAssess)
Manly.plot()	demo(DensityPlot), demo(ContourPlot)
Comprehensive analysis	demo(utility), demo(ais), demo(seeds), demo(bankruptcy) demo(acidity)

Table 3: Summary of demo examples included in **ManlyMix**.

Summary

The R package **ManlyMix** is discussed and illustrated in detail. The provided functions enable practitioners to analyze heterogeneous data and conduct cluster analysis with Manly mixture models. The algorithms behind functions are introduced and explained carefully. Illustrative examples based on challenging real-life datasets are studied to demonstrate the usefulness and efficiency of the package. Promising results suggest that **ManlyMix** is not only a powerful package for clustering and classification, but also a diagnostic tool to investigate skewness and deviation from normality in data. Demo examples are provided for each function in **ManlyMix** for the users to study.

Demo examples

For users who need further information about the package, we have constructed 16 demo examples listed in Table 4 that provide a comprehensive demonstration of **ManlyMix** capabilities. Among the examples, 11 of them are designed to demonstrate the capability and utility of each function and 5 of them run comprehensive analysis of classification datasets. Each demo can be accessed by its name and the users can reproduce themselves.

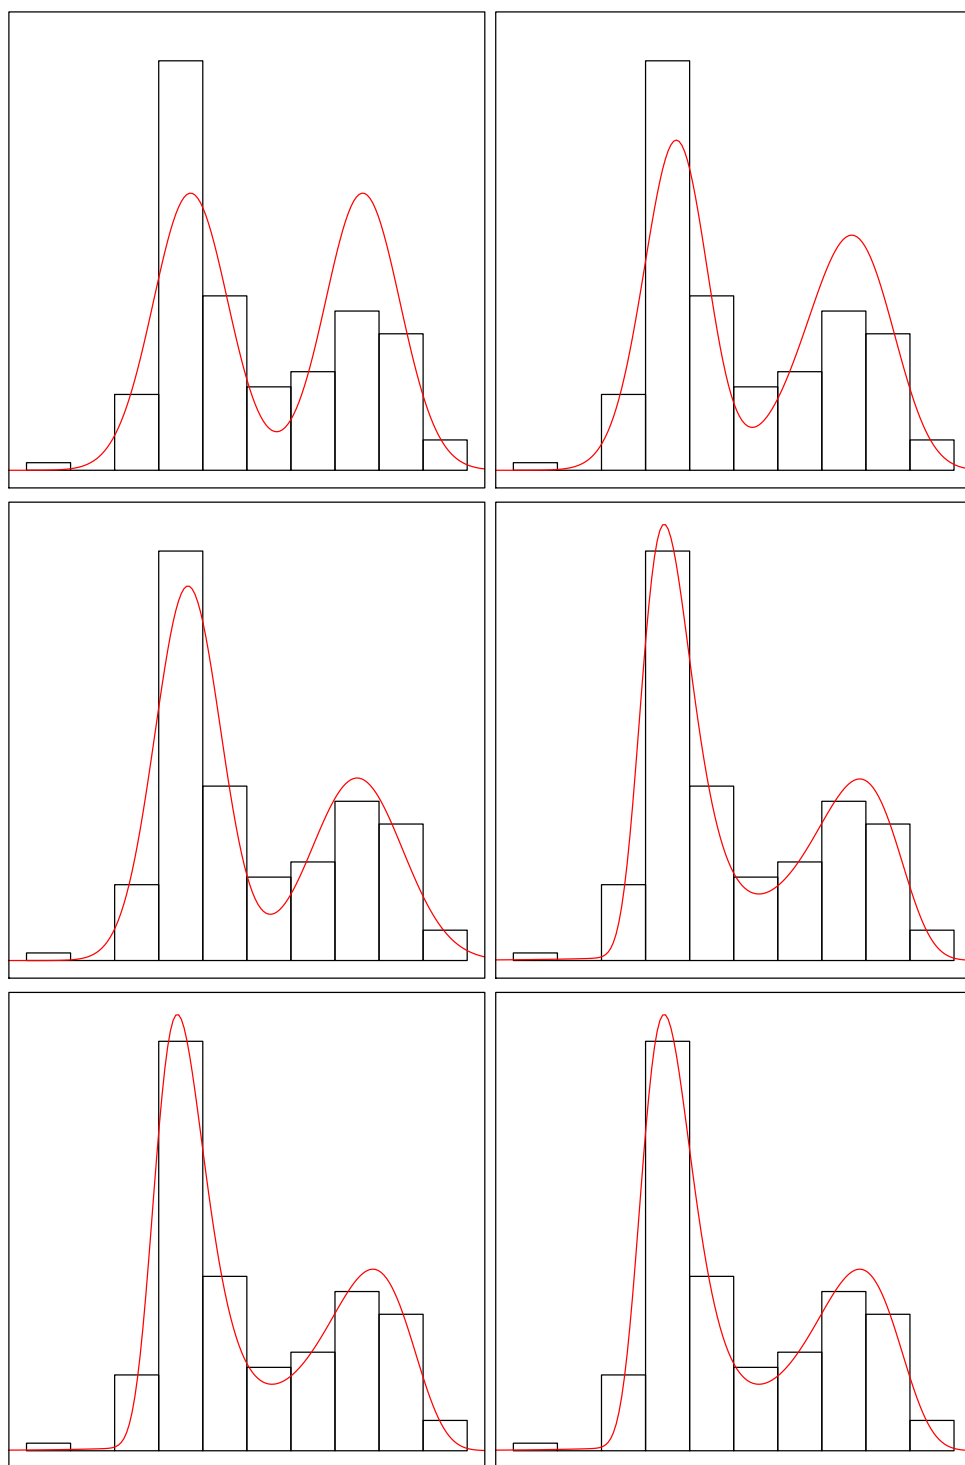


Figure 2: *acidity* dataset: fitted density plots from function `Manly.plot()`: K-means (first row left), Manly K-means (first row right), Gaussian mixture (second row left), Manly mixture (second row right), Manly forward (third row left) and Manly backward (third row right).

As an illustration of how these demos can be employed, the code of the first example can be approached through running the following code in R.

```
library(ManlyMix)
demo(EMalgorithm1)
```

Function	Demo example(s)
Manly.EM()	demo(EMalgorithm1), demo(EMalgorithm2)
Manly.select()	demo(ForwardSelection), demo(BackwardSelection)
Manly.Kmeans()	demo(ManlyKmeans1), demo(ManlyKmeans2)
Manly.overlap()	demo(Overlap)
Manly.sim()	demo(DataSimulation)
Manly.var()	demo(VarAssess)
Manly.plot()	demo(DensityPlot), demo(ContourPlot)
Comprehensive analysis	demo(utility), demo(ais), demo(seeds), demo(bankruptcy) demo(acidity)

Table 4: Summary of demo examples included in **ManlyMix**.

Bibliography

- E. Anderson. The Irises of the Gaspé peninsula. *Bulletin of the American Iris Society*, 59:2–5, 1935. [p6, 8]
- C. Biernacki, G. Celeux, and G. Govaert. Choosing starting values for the EM algorithm for getting the highest likelihood in multivariate Gaussian mixture models. *Computational Statistics & Data Analysis*, 413:561–575, 2003. [p16]
- G. Celeux and G. Govaert. A classification EM algorithm for clustering and two stochastic versions. *Computational Statistics & Data Analysis*, 14:315–332, 1992. URL [https://doi.org/10.1016/0167-9473\(92\)90042-e](https://doi.org/10.1016/0167-9473(92)90042-e). [p5]
- D. Cook and S. Weisberg. *An Introduction to Regression Graphics*. John Wiley & Sons, 1994. [p6, 12]
- S. L. Crawford. An application of the laplace method to finite mixture distribution. *Journal of the American Statistical Association*, pages 259–267, 1994. URL <https://doi.org/10.1080/01621459.1994.10476467>. [p17]
- A. P. Dempster, N. M. Laird, and D. B. Rubin. Maximum likelihood for incomplete data via the EM algorithm (with discussion). *Journal of the Royal Statistical Society B*, 39:1–38, 1977. [p2]
- R. A. Fisher. The use of multiple measurements in taxonomic problems. *The Annals of Eugenics*, 7: 179–188, 1936. URL <https://doi.org/10.1111/j.1469-1809.1936.tb02137.x>. [p6, 8]
- N. A. Gillespie and M. C. Neale. A finite mixture model for genotype and environment interactions: Detecting latent population heterogeneity. *Twin Res Hum Genet*, 9(3):412–23, 2006. URL <https://doi.org/10.1375/18324270677591380>. [p1]
- S. X. Lee and G. J. McLachlan. Model-based clustering and classification with non-normal mixture distributions. *Statistical Methods and Applications*, 22(4):427–454, 2013. [p12, 14]
- S. X. Lee and G. J. McLachlan. **EMMIXuskew**: Fitting unrestricted multivariate skew t mixture models. *Journal of Statistical Software*, 55(12), 2014. URL <https://doi.org/10.1007/s11222-012-9362-4>. [p1]
- K. Lo, F. Hahne, R. R. Brinkman, and R. Gottardo. **flowClust**: a bioconductor package for automated gating of flow cytometry data. *BMC Bioinformatics*, 10(145), 2009. URL <https://doi.org/10.1186/1471-2105-10-145>. [p1]
- R. Maitra and V. Melnykov. Simulating data to study performance of finite mixture modeling and clustering algorithms. *Journal of Computational and Graphical Statistics*, 19(2):354–376, 2010. URL <https://doi.org/10.1198/jcgs.2009.08054>. [p3]
- B. F. J. Manly. Exponential data transformations. *Biometrics Unit*, 25:37–42, 1976. [p2]
- G. J. McLachlan and K. E. Basford. *Mixture Models: Inference and Applications to Clustering*. Marcel Dekker, New York, 1988. URL <https://doi.org/10.2307/2348072>. [p3]

- G. J. McLachlan and T. Krishnan. *The EM Algorithm and Extensions*. John Wiley & Sons, New York, 2nd edition, 2008. URL <https://doi.org/10.1002/9780470191613>. [p2]
- G. J. McLachlan and D. Peel. *Finite Mixture Models*. John Wiley & Sons, New York, 2000. [p4, 9]
- V. Melnykov. On the distribution of posterior probabilities in finite mixture models with application in clustering. *Journal of Multivariate Analysis*, 122:175–189, 2013. URL <https://doi.org/10.1016/j.jmva.2013.07.014>. [p4]
- S. Michael and V. Melnykov. Finite mixture modeling of gaussian regression time series with application to dendrochronology. *Journal of Classification*, 33(3):412–441, 2016. URL <https://doi.org/10.1007/s00357-016-9216-4>. [p1]
- B.-J. Park and D. Lord. Application of finite mixture models for vehicle crash data analysis. *Accident Analysis & Prevention*, 41(4):683–691, 2009. URL <https://doi.org/10.1016/j.aap.2009.03.007>. [p1]
- M. O. Prates, C. R. B. Cabral, and V. H. Lachos. **mixsmsn**: Fitting finite mixture of scale mixture of skew-normal distributions. *Journal of Statistical Software*, 54:1–20, 2013. URL <https://doi.org/10.18637/jss.v054.i12>. [p1]
- P. Schlattmann. *Medical Applications of Finite Mixture Models*. Springer-Verlag, 2009. [p1]
- G. Schwarz. Estimating the dimensions of a model. *The Annals of Statistics*, 6:461–464, 1978. [p4]
- K. Wang, A. Ng, and G. McLachlan. **EMMIXskew**: The EM algorithm and skew mixture distribution, 2013. URL <https://doi.org/10.1109/dicta.2009.88>. R package version 1.0.1. [p1]
- J. H. Ward. Hierarchical grouping to optimize an objective function. *Journal of the American Statistical Association*, 58:236–244, 1963. [p11]
- X. Zhu and V. Melnykov. Manly transformation in finite mixture modeling. *Computational Statistics & Data Analysis*, 2016a. URL <https://doi.org/10.1016/j.csda.2016.01.015>. [p1]
- X. Zhu and V. Melnykov. **ManlyMix**: An R package for model-based clustering with manly mixture models, 2016b. R package version 0.1.9. [p1]

Xuwen Zhu

Department of Mathematics, The University of Louisville
Louisville, KY 40208
USA
xuwen.zhu@louisville.edu

Volodymyr Melnykov

Department of Information Systems, Statistics, and Management Science, The University of Alabama
Tuscaloosa, AL 35487
USA
vmelnykov@cba.ua.edu