

# cchs: An R Package for Stratified Case–Cohort Studies

by Edmund Jones

**Abstract** The `cchs` package contains a function, also called `cchs`, for analyzing data from a stratified case–cohort study, as used in epidemiology. For data from a stratified case–cohort study, `cchs` calculates Estimator III of Ørnulf Borgan et al. (2000), which is a score-unbiased estimator for the regression coefficients in the Cox proportional hazards model. From the user’s point of view the function is similar to `coxph` and other widely used model-fitting functions. Convenient software was not previously available for Estimator III because it is complicated to calculate. SAS and S-Plus code-fragments for the calculation have been published (Langholz and Jiao, 2007a; Cologne et al., 2012), but `cchs` is much easier to use on new datasets, and works differently so that it is far more efficient in terms of time and memory and can cope with much bigger datasets. It also avoids several minor approximations and simplifications.

## Introduction

One common type of study in epidemiology is the cohort study, in which participants are recruited and then followed up over time. Participants who experience a certain event, for example a heart attack or the diagnosis of a disease, are called cases, and participants who do not are called non-cases or controls. A related type of study is the case–cohort study, which is nested within a cohort study. In a case–cohort study, some of the covariates are measured not for the whole cohort but only for the cases and the subcohort, which is a randomly selected subset of the cohort. The advantage is that a case–cohort study can often estimate effect-sizes almost as accurately as the corresponding cohort study but at much lower cost, because measuring all the covariates on all the participants is expensive—the full cohort study would waste money by measuring the covariates on an excessive number of non-cases.

The case–cohort dataset contains only the cases and the subcohort, and is usually analyzed with a Cox proportional hazards regression model (Cox, 1972). The Cox model is often fitted using the estimator of Prentice (1986), who proposed the case–cohort design. More advanced kinds of analysis sometimes use information on the full cohort, such as the values of the covariates that were cheap to measure.

In a stratified case–cohort study, the selection of the subcohort is stratified. The cohort is divided into strata and each stratum is assigned a sampling fraction, which is the number of participants in that stratum who will be in the subcohort divided by the total number of participants in the stratum. Within a stratum, each participant is equally likely to be in the subcohort (but across the whole cohort the same is not true, unless all the sampling fractions are equal). The purpose of the stratification is usually to improve the efficiency of the estimator. To fit a Cox model to data from a stratified case–cohort study, one possible estimator is the time-fixed version of Estimator III of Ørnulf Borgan et al. (2000)—hereafter called “Borgan’s estimator”. Borgan’s estimator is optimal in the sense that it is score-unbiased; this criterion is explained in the section “Comparison with other estimators”.

The subject of this article is the `cchs` package (version 0.4.0; Jones, 2017), which is available on CRAN and has made Borgan’s estimator available in a convenient form for the first time. The package consists of a single function, also called `cchs`, which calculates the estimator, and a dataset on which the function can be used, called `cchsData`. The name `cchs` stands for “case–cohort stratified”.

Previous software for calculating Borgan’s estimator includes SAS code (SAS Institute Inc., 1999) that appears in Langholz and Jiao (2007a), and S-Plus code (Tibco Software Inc., 2007) in the appendix of Cologne et al. (2012). S-Plus is no longer available to buy but it is very similar to R; Cologne et al. (2012)’s code works in R if one function, `match.data.frame`, is imported from S-Plus or rewritten (it contains 15 lines of code, of which 8 check the arguments).

The previously published SAS code and S-Plus code are both for analyzing specific datasets with specific models. They could be rewritten to work more generally, but even then with some datasets they would give slightly inaccurate results and with others they would use far too much computational time or memory, as explained in subsequent sections. (Langholz and Jiao wrote more general SAS code, available at Langholz and Jiao, 2007b, but this still has the problems with inaccurate results and computational resources and is less easy to use than `cchs`.) Cologne et al. (2012) state that Borgan’s estimator can be calculated by an application called *Epicure*, which is now sold by Risk Sciences International. The estimator was also calculated by Borgan et al. themselves, but their code is not publicly available.

The `cch` function, in the [survival](#) package (Therneau, 2017; Therneau and Grambsch, 2000), can calculate several estimators for case-cohort data, including the estimator of Prentice (1986) and Estimators I and II of Ørnulf Borgan et al. (2000). So `cch` corresponds to just one feature that could be added to `cch` as an extra option. But Estimator III is more complicated to calculate than the other estimators and `cch` has numerous features that `cch` lacks. Other software for stratified case-cohort studies includes the [survey](#) and [NestedCohort](#) packages, which use inverse probability weights (Lumley, 2004, 2017a; Mark and Katki, 2006; Katki and Mark, 2008), and the code for Estimator II in Samuelsen et al. (2007).

The next five sections of the article are about the model, the estimator, and how to use `cch`. The function is easy to use for anyone who is accustomed to model-fitting functions such as `lm` or `coxph`. After that are several sections about how `cch` works internally and how the previously published SAS and S-Plus code-fragments work. The previously published code makes several approximations and is unable to cope with large datasets, but `cch` is implemented differently and avoids all these drawbacks. `cch` manipulates the data in an economical way, leading to huge savings in computational time and memory. The novel implementation illustrates how such savings can sometimes be achieved by high-level data manipulations as opposed to lower-level methods such as optimization algorithms or parallelization.

Parts of `cch` are based on code from `coxph`, `cch`, and the appendix of Cologne et al. (2012), and the formulas and some of the computational methods are from Therneau and Li (1999), Ørnulf Borgan et al. (2000), Langholz and Jiao (2007a), and Cologne et al. (2012). But a major part of the computational methods is original (see the section “How `cch` works”), and several other issues are expounded in this article for the first time (see “Obscure aspects of the calculation”).

The idea of creating `cch` came from work on EPIC-InterAct (InterAct Consortium, 2011) and EPIC-CVD (Danesh et al., 2007), two stratified case-cohort studies that are nested within the cohort from EPIC, the European Prospective Investigation of Cancer and Nutrition (Bingham and Riboli, 2004). EPIC-InterAct is a study of incident type 2 diabetes with data from 29 centres in eight European countries. The centres range from Ragusa in Sicily to Umea in northern Sweden. The selection of the subcohort was stratified by centre or country. EPIC-CVD is a study of cardiovascular disease that uses most of the same centres as EPIC-InterAct and the same subcohort in those centres. In Jones et al. (2015), data from EPIC-InterAct was used to compare five different models for data from stratified case-cohort studies, and a pre-release version of `cch` was used to fit two of those models.

## The model and the estimator

The Cox model for survival-time data is commonly defined by this equation:

$$h_i(t) = h_0(t) \exp(\beta^\top z_i),$$

where

$h_i$  is the hazard function for participant  $i$

$t$  is the time to the event

$h_0$  is the baseline hazard function

$\beta$  is the regression coefficients (which are log hazard ratios)

$z_i$  is the covariates for participant  $i$ .

If all the event-times are distinct then the partial likelihood, as used with data from a cohort study or clinical trial, is

$$\prod_j \frac{\exp(\beta^\top z_{i_j})}{\sum_k Y_k(t_j) \exp(\beta^\top z_k)},$$

where

$t_j$  is the  $j$ th event-time

$i_j$  is the participant whose event happened at time  $t_j$

$$Y_k(t_j) = \begin{cases} 1 & \text{if } t_j \in (t_{0k}, t_{1k}] \\ 0 & \text{otherwise} \end{cases}$$

$t_{0k}$  is the entry-time for participant  $k$  (this is often zero for all participants)

$t_{1k}$  is the exit-time for participant  $k$  (the time at which they had the event or were censored).

The product is over all the cases and the sum is over all the participants. The time-interval  $(t_{0k}, t_{1k}]$  is participant  $k$ 's "at-risk time", so  $Y_k(t_j)$  is an indicator variable for whether participant  $k$  is at risk at time  $t_j$ .

For case-cohort data,  $z_i$  is only known for the cases and subcohort members, so the partial likelihood cannot be evaluated and it is impossible to estimate the regression coefficients by the method of maximum likelihood. Estimators are therefore defined in terms of pseudo-likelihoods, which are approximations of the partial likelihood. Each estimator is the value of  $\beta$  that maximizes the corresponding pseudo-likelihood. The pseudo-likelihood for Borgan's estimator is

$$\prod_j \frac{\alpha_{s(i_j)}^{-1} \exp(\beta^\top z_{i_j})}{\sum_{k \in R(t_j)} Y_k(t_j) \alpha_{s(k)}^{-1} \exp(\beta^\top z_k)},$$

where

$s(i)$  is the stratum that contains participant  $i$

$\alpha_s$  is the sampling fraction for stratum  $s$  (the proportion of stratum  $s$  that appears in the subcohort)

$$R(t_j) = \begin{cases} C & \text{if } i_j \in C \\ C \cup \{i_j\} \setminus \{J_{s(i_j)}\} & \text{if } i_j \notin C \end{cases}$$

$C$  is the subcohort

$J_s$  is a participant randomly selected from  $C \cap s$  (where  $s$  is a stratum).

Ørnulf Borgan et al. (2000) called this a "swapper" method because of the way  $R(t_j)$  is defined for  $i_j \notin C$ :  $J_{s(i_j)}$  is removed from  $C$  and  $i_j$  is added. If  $J_{s(i_j)}$  is not at risk at time  $t_j$  then its removal from  $R(t_j)$  makes no difference to the pseudo-likelihood, because  $Y_{J_{s(i_j)}}(t_j) = 0$ .

For Borgan's estimator it is recommended to use an asymptotic covariance estimator rather than a robust one (Jiao, 2001). This is explained in "The calculation of the covariance matrix".

Langholz and Jiao (2007a) discuss two situations in which a case-cohort study might be stratified, which they call "exposure stratified" and "confounder stratified". In the exposure stratified situation, the strata are defined by levels of one or more covariates, and the purpose of the stratification is to increase the efficiency of the estimator. In the confounder stratified situation, the strata are defined by a covariate that is believed to confound the relation between the exposure of interest and the event-time, for example the centre in a multi-centre study. For an exposure stratified study, it is usual to fit a single unstratified Cox model to the data, and this is what Borgan's estimator is designed for. For a confounder stratified study, it is more natural to fit a stratified Cox model—in which each stratum  $s$  has its own baseline hazard  $h_{0s}(t)$  and  $h_i(t) = h_{0s(i)}(t) \exp(\beta^\top z_i)$ —and for this model the estimator of Prentice (1986) should be used; but it is still possible to fit an unstratified model using Borgan's estimator.

## Comparison with other estimators

Borgan's estimator is "score-unbiased." This is one of the main desirable criteria for estimators for case-cohort data. It means that the expectation of the pseudo-score (the derivative of the pseudo-likelihood) is zero when the coefficients take their true values. See Godambe (1960, 1976), Lindsay (1982), Severini (2011), or Cologne et al. (2012). The removal of the randomly selected element  $J_{s(i_j)}$ , in the definition of  $R(t_j)$ , is done with the explicit purpose of adjusting the pseudo-score so that its expectation is zero.

Ørnulf Borgan et al. (2000) described three estimators for the Cox model with data from a stratified case-cohort study. These have been investigated in several simulation studies: Borgan et al. used cohorts of size 1000 and subcohorts of size 100 and found that there was little to choose between Estimators II and III, but both performed better than Estimator I; Cologne et al. (2012) made similar findings but reported that Estimator III did slightly better than Estimator II when the sampling fraction was very small; and Samuelsen et al. (2007) found that Estimator II performed well. As for theoretical properties, of the three only Estimator III is score-unbiased, but Estimator II has smaller asymptotic variance (this was hinted at in Ørnulf Borgan et al., 2000, stated in Samuelsen et al., 2007, and explained most clearly in Samuelsen et al., 2006).

Borgan et al. also described time-dependent versions of the three estimators; these have smaller asymptotic variance (Kulich and Lin, 2004), and the time-dependent version of Estimator III is score-unbiased (Ørnulf Borgan et al., 2000), but it would be even more complicated to calculate than the time-fixed version calculated by cchs. Estimator II can be improved using covariate data on

the whole cohort, if that is available (Breslow et al., 2009a,b; Yan et al., 2017). Ørnulf Borgan et al. (2000) mentioned that Prentice (1986)'s estimator for unstratified case-cohort studies, for which the pseudo-likelihood is

$$\prod_j \frac{\exp(\beta^\top z_{i_j})}{\sum_{k \in C \cup \{i_j\}} Y_k(t_j) \exp(\beta^\top z_k)},$$

can be adapted to work with stratified studies, but is only score-unbiased if for  $i_j \notin C$  the  $C$  in the above expression is replaced by  $(C \cap s(i_j))$ . This is less efficient than Estimator III, because for these  $i_j$  the denominator includes fewer participants.

Cox models for stratified case-cohort data can also be fitted with general methods that use inverse probability weights. These methods are suitable for a wide range of complex sampling schemes (Mark and Katki, 2006; Lumley, 2017b), but no claim is made that the estimators are score-unbiased. This approach was used by Gray (2009), Liu et al. (2012), and Payne et al. (2016). In other research, Kulich and Lin (2004) presented a framework that covered many case-cohort estimators and the possibility of stratified subcohort-selection, and Zeng and Lin (2007) and Kim et al. (2013) described wide classes of sampling schemes and regression models for survival analysis, of which the stratified case-cohort design and the Cox model are special cases.

## How to use cchs

From the user's point of view cchs is reasonably similar to coxph. The most important arguments are:

formula, which specifies the model; the left-hand side must be a "Surv" object  
data, a data-frame or environment that contains the variables in the formula  
inSubcohort, an indicator for whether each participant is in the subcohort  
stratum, the stratum variable  
samplingFractions, the sampling fraction for each stratum  
cohortStratumSizes, the size of each stratum in the full cohort  
precision (see the text).

If the data is supplied as a data-frame, it is assumed to be in the usual form where each row corresponds to one observation (that is, one participant in the case-cohort study) and each column corresponds to a variable. inSubcohort and stratum can either be in data or be separate objects that are named explicitly in the call to cchs.

Either samplingFractions or cohortStratumSizes must be given, but not both. If the latter is given then then cchs uses it to calculate the former—for each stratum, it divides the size of the subcohort, which can be found from the data, by the size of the cohort, which is given by cohortStratumSizes. If the sampling fractions are not stored as a column of the data then they can be passed to cchs as a named vector. For example, if the stratum variable takes the values "London" and "Paris" and the sampling fractions are 0.1 and 0.2 respectively, then this vector should be `c(London = 0.1, Paris = 0.2)`.

The precision argument must be set if there are any tied event-times. Borgan's estimator requires all event-times to be distinct, so cchs changes tied event-times by small amounts, as suggested by Langholz and Jiao (2007a), and it uses precision to work out suitable values for these amounts. If the times are all integers then precision should be set to 1. In datasets from EPIC-InterAct and EPIC-CVD, the times are recorded to the nearest day but stored as numbers of years, so precision needs to be set to `1 / 365.25`.

cchs has several other arguments, which should normally be left to take their default values. Several of these are discussed later, in "Obscure aspects of the calculation".

## An example analysis

cchsData is an artificial stratified case-cohort dataset that was created from the nwtco dataset in the survival package. nwtco comes from two clinical trials run by the National Wilms Tumor Study Group in the United States (D'Angio et al., 1989; Green et al., 1998). Different artificial case-cohort datasets based on nwtco were used by Kulich and Lin (2004), Breslow et al. (2009a), and Breslow et al. (2009b).

In cchsData the event is relapse of Wilms' tumour and the subcohort-selection is stratified by the result of a histological examination. The most important variables are time (the time to the event or censoring), isCase (the event-indicator), inSubcohort (the indicator for being in the subcohort),

localHistol (the stratum variable), and sampFrac (the stratum-specific sampling fractions). The times are stored as numbers of days, and some of them are tied, so the precision argument of cchs has to be set to 1 (or 1/2, 1/3, etc.).

The sampling fractions are 5% and 20% and the subcohort is about half as big as the set of cases. In case-cohort studies in general, the sampling fraction is typically in the area of 5% but sometimes larger (Sharp et al., 2014), and the subcohort is usually larger than the set of cases (Juraschek et al., 2013; Lamb et al., 2013; Jones et al., 2015) but occasionally smaller (Huxley et al., 2013).

The following command uses Borgan's estimator to fit a Cox model to cchsData in which the covariates are age at diagnosis (ageAtDiagnosis) and stage of cancer at diagnosis (stage):

```
> cchs(Surv(time, isCase) ~ ageAtDiagnosis + stage, data = cchsData,
+ inSubcohort = inSubcohort, stratum = localHistol,
+ samplingFractions = sampFrac, precision = 1)
```

cchs returns an object of S3 class "cchs". This contains most of the same elements as a "coxph" object and several extra elements. When a "cchs" object is printed, information about the dataset and the changes to the tied event-times appears before the coefficients. The output from the above command is:

```
Call: cchs(formula = Surv(time, isCase) ~ stage + ageAtDiagnosis,
  data = cchsData, inSubcohort = inSubcohort, stratum = localHistol,
  samplingFractions = sampFrac, precision = 1)
Number of observations/rows: 785, of which
  subcohort non-cases: 214
  subcohort cases:    48
  non-subcohort cases: 523
179 of 571 discretized event-times were changed by up to 0.01 to deal
  with ties.
Number of strata for subcohort-selection: 2
Coefficients (HR = hazard ratio, CI = 95% confidence interval, SE = standard
error):
```

	HR	CIlower	CIupper	p	logHR	SElogHR
ageAtDiagnosis	1.101685	1.029489	1.178944	0.005104173	0.09684087	0.03458127
stageII	1.397695	0.915168	2.134635	0.121219767	0.33482413	0.21606100
stageIII	1.794611	1.156081	2.785816	0.009150320	0.58478850	0.22436755
stageIV	2.367953	1.382691	4.055282	0.001686931	0.86202589	0.27449190

This can be interpreted like output from any other function that fits the Cox model. For example, the hazard is estimated to be 79% higher for patients who were at stage III when they were diagnosed than for patients who were at stage I, adjusting for age at diagnosis, and the 95% confidence interval for this quantity is (16%, 179%). The level for the confidence intervals in the output is determined by the confidenceLevel argument to cchs, for which the default is 0.95.

If the output of cchs is stored as an object called result then the coefficients are available as result\$coefficients and the covariance matrix is available as result\$var. The table at the bottom of the output is stored as result\$coeffsTable, so that the confidence intervals for the hazard ratios and the other quantities in that table are easily available. For example, to get the confidence intervals for the hazard ratio for ageAtDiagnosis, type

```
> result <- cchs(Surv(time, isCase) ~ ageAtDiagnosis + stage, data = cchsData,
+ inSubcohort = inSubcohort, stratum = localHistol,
+ samplingFractions = sampFrac, precision = 1)
> result$coeffsTable["ageAtDiagnosis", c("CIlower", "CIupper")]
```


## Other issues for the user

Calling cchs twice with the same arguments will not necessarily give exactly the same results. This is obviously unusual for estimators that maximize a likelihood or pseudo-likelihood. It happens firstly because Borgan's estimator involves randomness in the choice of  $J_{s(i_j)}$  and secondly because if there are tied event-times then those are changed in a random way.

To get exactly the same results every time, set the random seed (using set.seed) just before the call to cchs. Another approach would be to call cchs multiple times, using the replicate command. One possible way to then combine the results might be Rubin's rules (Rubin, 1987; Schafer, 1999), as used in multiple imputation. For each term in the model, the coefficient estimate is simply the mean



of the separate estimates, and the variance is the mean of the separate variances plus a term to account for the between-replication variability. It might be worth expressly comparing the within-replication variability (the separate variance estimates) with the between-replication variability (the variability due to the randomness of the estimator). As an example, Figure 1 shows the point-estimates and confidence intervals that appear in the output in the previous section, and for comparison the quartiles and ranges of the point-estimates produced by 1000 calls to `cchs` with the same arguments. Here the between-replication variability is small, and very small compared to the within-replication variability.



**Figure 1:** The dots and horizontal lines show the point-estimates and confidence intervals from the output in “An example analysis”; this is within-replication variability. The boxes and whiskers show the hazard ratios from 1000 replications of `cchs`, using the dataset and model in the same section; this is between-replication variability. The boxes show the quartiles and the whiskers show the full ranges.

In `coxph`, model formulas can contain certain special terms. Of these, `cchs` allows offset terms, but it does not allow `cluster` (for identifying correlated groups of rows), `strata` (for fitting a stratified Cox model), or `tt` (for time-varying covariates). If the data are from a stratified case-cohort study and it is desired to fit a stratified Cox model where the strata in the baseline hazard are the same as the strata for the subcohort-selection, then Prentice’s estimator should be used (Langholz and Jiao, 2007a).

`cchs` does a complete-case analysis. Participants who have missing values (NAs) for `inSubcohort`, `stratum`, or any of the variables that appear in the model-formula (including the “Surv” object) are dropped.

All error and warning messages produced by `cchs` are designed to be meaningful. This is achieved by checking the arguments and raising an error if any of them have illegal values, and by adding extra messages to any errors or warnings raised by `model.frame` or `coxph` when they are called internally by `cchs`. The addition of the extra messages results in the call stack, which can be viewed by typing `traceback()`, being longer and more complicated than usual. So if `traceback()` is going to be used to diagnose problems, it is advisable to switch the extra messages off by setting `annotateErrors = FALSE`.

## How `cchs` works

For users it should not be necessary to understand the internal workings of `cchs`. But `cchs` works differently from the previously published SAS and S-Plus code in several ways, and these novel aspects of its implementation affect the amount of computational time and memory that it uses and enable it to work with much bigger datasets. These aspects also affect the numbers in the output. This section and the next three explain the details. They should be helpful to cautious or curious users who want to know exactly how `cchs` works and to anyone who intends to compare it with other software for Borgan’s estimator.

Functions that calculate estimators for the Cox model with case-cohort data invariably work by manipulating the data or the model-formula in various ways—for example, duplicating rows and changing entry-times, or appending weights to the rows and including the weight in the model-formula—and then passing those to a second function (in R, `coxph`) that was originally designed to fit the Cox model with the standard estimator. The manipulation is done in such a way that when the

second function attempts to calculate and maximize the partial likelihood, what it actually calculates and maximizes is the pseudo-likelihood for the case-cohort estimator. The second function returns the estimates of the regression coefficients and their covariance matrix, and the final step is usually to make adjustments to the covariance matrix (see the next section).

The reason why this method works is that the standard partial likelihood and the various pseudo-likelihoods all have similar forms. They have the same number of multiplicative terms, one for each case, and the same values in the numerators (apart from multiplicative factors, which are easy to deal with). The only differences are the sets of participants who appear in the denominators and, for some pseudo-likelihoods, multiplicative factors in the denominators. These sets are dealt with by duplicating or excluding rows and changing entry- and exit-times, and any multiplicative factors are dealt with by using an offset term in the model-formula.

For Borgan's estimator the manipulations are rather complicated because of the need to deal with the "swapping". `cchs` manipulates the data as follows:

1. Define a small number  $\epsilon$ , which is smaller than any of the differences between consecutive event-times.
2. For each  $i_j \notin C$ , change the entry-time to  $t_j - \epsilon$  (where  $t_j$  is the exit-time).
3. For each stratum  $s$ :
  - 3.1 Choose a random participant  $J_s$  from  $C \cap s$ . Denote the entry-time for  $J_s$  by  $t_{0s}$  and the exit-time by  $t_{1s}$ .
  - 3.2 Let  $U_s = \{t_j : i_j \in s, i_j \notin C, t_{0s} < t_j \leq t_{1s}\}$ . Sort  $U_s$  into increasing order as  $u_{s(1)}, \dots, u_{s(|U_s|)}$ .
  - 3.3 Replace  $J_s$ 's row by  $|U_s| + 1$  new rows with variables as follows.
    - Row 1: entry-time  $t_{0s}$ , exit-time  $u_{s(1)} - \epsilon$ , event-indicator zero.
    - Row 2: entry-time  $u_{s(1)}$ , exit-time  $u_{s(2)} - \epsilon$ , event-indicator zero.
    - ...
    - Row  $|U_s|$ : entry-time  $u_{s(|U_s|-1)}$ , exit-time  $u_{s(|U_s|)} - \epsilon$ , event-indicator zero.
    - Row  $|U_s| + 1$ : entry-time  $u_{s(|U_s|)}$ , exit-time  $t_{1s}$ , event-indicator as in the original row.

These steps are best understood by referring to the earlier definition of  $R(t_j)$  and bearing in mind that `coxph` takes a row's at-risk time to be the half-open interval (entry-time, exit-time]. Step 1 will be explained last. Step 2 has the effect of removing the non-subcohort cases from  $R(t_j)$ , except for  $i_j$  itself, which is not removed.

Step 3 makes the further changes to  $R(t_j)$  that are needed if  $i_j \notin C$ . Note that it deals with one  $J_s$ , not one  $i_j$ , at a time. Step 3.1 chooses  $J_s$ . Step 3.2 defines  $U_s$  to be the event-times that lie in the range of  $J_s$ 's at-risk time and correspond to non-subcohort cases in  $s$ . Step 3.3 excises a short stretch of at-risk time from  $J_s$  at each time in  $U_s$  by splitting  $J_s$ 's row into multiple rows with separate stretches of at-risk time.

The effect of this is to remove  $J_{s(i_j)}$  from  $R(t_j)$  when  $i_j \notin C$ , as required by the definition of  $R(t_j)$ . For a non-subcohort case in  $s$  whose event-time is not in  $U_s$ , it makes no difference whether  $J_{s(i_j)}$  is in  $R(t_j)$  or not, because  $Y_{J_{s(i_j)}}(t_j)$  is zero; so for such cases this manipulation is not necessary.

The purpose of step 1 should now be clear. It defines  $\epsilon$  to be sufficiently small that steps 2 and 3 do not have any unintended effects on other terms in the pseudo-likelihood.

One other change to the data is needed. An extra column is created that contains  $-\log \alpha_{s(k)}$  for participant  $k$ , and this is included as an offset term in the model-formula that is passed to `coxph`. This has the effect of inserting the  $\alpha_{s(k)}^{-1}$  factors that are needed in the denominators.

## The calculation of the covariance matrix

`cchs` uses an asymptotic covariance estimator, as recommended by Jiao (2001). Internally, after it manipulates the data and calls `coxph`, the final step is to adjust the covariance matrix that is returned by `coxph`. This is done by using the matrix of "dfbeta" residuals,  $D$ , which is produced by `residuals.coxph`;  $d_{ij}$  is the change in the estimate of regression coefficient  $j$  when row  $i$  is dropped from the data (or an approximation of that change). The covariance matrix is adjusted by adding  $\sum_s m_s(1 - \alpha_s) \text{Cov } D_{C \cap s}$ , where  $m_s = |C \cap s|$ ,  $D_{C \cap s}$  is the rows of  $D$  that correspond to  $C \cap s$ , and  $\text{Cov } D_{C \cap s}$  is the empirical covariance matrix  $D_{C \cap s}^T D_{C \cap s} / (m_s - 1)$ .

This adjustment corresponds to equation (17) in Ørnulf Borgan et al. (2000). The use of  $D$  is based on the formula for the adjustment to the covariance matrix with unstratified case-cohort data, which appears on page 102 of Therneau and Li (1999) and as equation (5) in Langholz and Jiao (2007a).

There is a small discrepancy between the formulas in these two publications. Let  $D_C$  be the rows of  $D$  that correspond to  $C$  and let  $\alpha = m/n$  be the sampling fraction ( $m$  is the size of the subcohort and  $n$  is the size of the cohort). According to Therneau and Li the quantity to add to the covariance matrix is  $(1 - \alpha)D_C^\top D_C$ , but according to Langholz and Jiao it is  $\frac{m(n-m)}{n} \text{Cov } D_C = (1 - \alpha) \frac{m}{m-1} D_C^\top D_C$ . (Langholz and Jiao use  $\text{Cov } D_C$  to mean  $D_C^\top D_C / (m - 1)$ , as can be seen from their Figures 2 and 4 and the SAS documentation for `proc corr`.) These differ by a factor of  $m / (m - 1)$ , which is obviously unimportant if  $m$  is large, but does make a difference if  $m$  is small. `cchs` does the calculation in the same way as Langholz and Jiao and [Cologne et al. \(2012\)](#).

## How the SAS code and S-Plus code work

The previously published SAS and S-Plus code-fragments for Borgan's estimator manipulate the data in a completely different way from `cchs`:

1. Define a small number  $\epsilon$ , which is smaller than any of the differences between consecutive event-times.
2. For each  $i_j \notin C$ , change the entry-time to  $t_j - \epsilon$ .
3. Replace the dataset with a new dataset that is constructed as follows. For each  $i_j$ :
  - 3.1 Make one copy of all the rows.
  - 3.2 Drop those of the new rows that are not at risk at  $t_j$ .
  - 3.3 For all rows, set the entry-time to  $t_j - \epsilon$  and the exit-time to  $t_j$ , and set the event-indicator to 1 for  $i_j$  and 0 for the other rows.
4. For each  $i_j \notin C$ , choose a random participant  $J_{s(i_j)}$  from  $C \cap s(i_j)$  and remove the row that contains  $J_{s(i_j)}$  and has exit-time  $t_j$ .

Step 3 radically changes the dataset and greatly increases its size. The result of the manipulations is that when the data is passed to `coxph` (in SAS, `proc phreg`), each row corresponds to exactly one additive term in the numerator or denominator of one multiplicative term of the pseudo-likelihood. The  $\alpha_{s(k)}^{-1}$  terms and the adjustment to the covariance matrix are dealt with in the same way as in `cchs`.

The advantage of this method is that steps 3 and 4 are relatively simple—the “swapping” in step 4 just consists of choosing and removing a single row each time. The disadvantages are that if the original dataset is medium or large in size then this method uses a gigantic amount of memory and requires `coxph` to deal with a gigantic dataset. For step 3.1 the S-Plus code creates a data-frame that contains one row for every combination of a case and a participant. In EPIC-CVD, a typical dataset that was supplied to an investigator contained about 31,000 participants, of whom 14,000 were cases, and this took up about 4.4MB as a binary file. So the data-frame created by step 3 would have  $31,000 \times 14,000 \approx 434$  million rows and take up about 62GB. This is too big for a desktop computer to store in memory, and even after some rows are removed in step 4 the data-frame would be far too big for `coxph` to process in a reasonable amount of time.

In contrast, when `cchs` is used on the same dataset, the manipulated dataset has fewer than  $31,000 + 14,000 = 45,000$  rows. This 45,000 was calculated by supposing that for every case an extra row is created, whereas in reality some cases are in the subcohort and it is likely that some “swapper” rows are not at risk at the relevant event-times, and for these no extra rows are created, so the true number of rows in the manipulated dataset is less than 45,000.

The manipulation of the data by the SAS and S-Plus code can be regarded as maximal, in the sense that each row ultimately corresponds to a single additive term in a single multiplicative term of the pseudo-likelihood. The manipulation by `cchs` can be regarded as minimal, since it makes only the manipulations that are strictly necessary in order for `coxph` to calculate the correct pseudo-likelihood, and it leaves the data with the smallest possible number of rows.

## Obscure aspects of the calculation

Three obscure issues arise in calculating Borgan's estimator or comparing `cchs` with the SAS and S-Plus code. Under normal circumstances users of `cchs` can ignore these—the logical arguments should be left to take their default values and the estimator will be calculated correctly. But for development or testing it may be necessary to understand these issues or set the logical arguments to different values.

The first issue is to do with the `dropNeverAtRiskRows` argument. This determines whether rows should be dropped if their at-risk time does not contain any of the event-times. (This only applies to



non-cases, since any case's at-risk time obviously contains that case's event-time.) Such rows make no difference to the regression coefficients that are output by `cchs`, but they do affect the covariance-estimates and confidence intervals, because residuals. `coxph` calculates the `dfbeta` residuals using an approximation, not exactly. The SAS code and S-Plus code drop these rows (in step 3.2), but Ørnulf Borgan et al. (2000) and Langholz and Jiao (2007a) make no mention of this issue. In `cchs`, `dropNeverAtRiskRows` is TRUE by default but can be set to FALSE if desired.

The second issue is to do with the `dropSubcohEventsDfbeta` argument. When the `dfbeta` residuals are used to calculate the adjustment to the covariance matrix, the correct way is to use all the rows of the matrix of `dfbeta` residuals that correspond to subcohort members. But the SAS code and S-Plus code both omit the rows that correspond to subcohort cases at their event-times (presumably because this makes the code simpler and is unlikely to change the output much; see "There is a slight approximation" in Section 2.4 of Langholz and Jiao, 2007a). By default, `cchs` calculates the adjustment to the covariance matrix in the strictly correct way. But if `dropSubcohEventsDfbeta` is TRUE, then it does it in the same way as the SAS and S-Plus code. To do this, for each  $i_j \in C$  it splits that row of the data at  $t_j - \epsilon$ , so that the first row has entry-time  $t_{0i_j}$ , exit-time  $t_j - \epsilon$ , and event-indicator 0 and the second row has entry-time  $t_j - \epsilon$ , exit-time  $t_j$ , and event-indicator 1; and when calculating the adjustment to the covariance matrix it excludes the row of the `dfbeta` residuals that corresponds to the second of these.

The third issue is that the "swapper"  $J_{s(i_j)}$  is supposed to be selected once for each stratum (see Section 3 of Ørnulf Borgan et al., 2000). `cchs` does this correctly, but the SAS and S-Plus both select  $J_{s(i_j)}$  once for each case. There is no practical way to change `cchs` to make the selection happen in the same way as in the SAS and S-Plus code, because `cchs` manipulates the data in a completely different way. So `cchs` does not have any argument to specify how the swappers should be selected.

## Discussion

In case-cohort studies, stratified selection of the subcohort seems to be common (Sharp et al., 2014), so there are likely to be plenty of investigations where `cchs` is useful. The way in which it manipulates the dataset makes it much faster and more capable than the previously published SAS and S-Plus code, and of course it is freely available. The three obscure issues described in the previous section are not greatly important, but they make a difference with small datasets and need to be considered when developing software to calculate the estimator.

Internally, `cchs` stores some redundant information, because when it splits a row it makes a copy of all the quantities in it, which might include a long list of covariates. However, `cchs` works with the model-matrix, not the original dataset, so covariates that are not in the model-formula are not duplicated. Storing a certain amount of redundant information is unavoidable given that the model-matrix is going to be passed to `coxph`—which is the only sensible way to calculate the estimator.

## Acknowledgments

The author would like to thank Michael Sweeting for his helpful comments. This work was funded by the UK Medical Research Council (G0800270), the British Heart Foundation (SP/09/002), the UK National Institute for Health Research (Cambridge Biomedical Research Centre), the European Research Council (268834), and the European Commission Framework Programme 7 (HEALTH-F2-2012-279233).

## Bibliography

- S. Bingham and E. Riboli. Diet and cancer—the European Prospective Investigation into Cancer and Nutrition. *Nature Reviews Cancer*, 4(3):206–215, 2004. URL <https://doi.org/10.1038/nrc1298>. [p2]
- N. E. Breslow, T. Lumley, C. M. Ballantyne, L. E. Chambless, and M. Kulich. Improved Horvitz–Thompson estimation of model parameters from two-phase stratified samples: Applications in epidemiology. *Statistics in Biosciences*, 1(1):32–49, 2009a. URL <https://doi.org/10.1007/s12561-009-9001-6>. [p4]
- N. E. Breslow, T. Lumley, C. M. Ballantyne, L. E. Chambless, and M. Kulich. Using the whole cohort in the analysis of case-cohort data. *American Journal of Epidemiology*, 169(11):1398–1405, 2009b. URL <https://doi.org/10.1093/aje/kwp055>. [p4]

- J. Cologne, D. L. Preston, K. Imai, M. Misumi, K. Yoshida, T. Hayashi, and K. Nakachi. Conventional case-cohort design and analysis for studies of interaction. *International Journal of Epidemiology*, 41(4): 1174–1186, 2012. URL <https://doi.org/10.1093/ije/dys102>. [p1, 2, 3, 8]
- D. R. Cox. Regression models and life-tables. *Journal of the Royal Statistical Society B*, 34(2):187–220, 1972. URL <https://www.jstor.org/stable/2985181>. [p1]
- J. Danesh, R. Saracci, G. Berglund, E. Feskens, K. Overvad, S. Panico, S. Thompson, A. Fournier, F. Clavel-Chapelon, M. Canonico, R. Kaaks, J. Linseisen, H. Boeing, T. Pischon, C. Weikert, A. Olsen, A. Tjønneland, S. P. Johnsen, M. K. Jensen, J. R. Quirós, C. A. G. Svatetz, M.-J. S. Pérez, N. Larrañaga, C. N. Sanchez, C. M. Iribas, S. Bingham, K.-T. Khaw, N. Wareham, T. Key, A. Roddam, A. Trichopoulos, V. Benetou, D. Trichopoulos, G. Masala, S. Sieri, R. Tumino, C. Sacerdote, A. Mattiello, W. M. M. Verschuren, H. B. B. de Mesquita, D. E. Grobbee, Y. T. van der Schouw, O. Melander, G. Hallmans, P. Wennberg, E. Lund, M. Kumle, G. Skeie, P. Ferrari, N. Slimani, T. Norat, and E. Riboli. EPIC-Heart: The cardiovascular component of a prospective study of nutritional, lifestyle and biological factors in 520,000 middle-aged participants from 10 European countries. *European Journal of Epidemiology*, 22(2):129–141, 2007. URL <https://doi.org/10.1007/s10654-006-9096-8>. [p2]
- G. J. D’Angio, N. Breslow, J. B. Beckwith, A. Evans, E. Baum, A. Delorimier, D. Fernbach, E. Hrabovsky, B. Jones, P. Kelalis, H. B. Othersen, M. Tefft, and P. R. M. Thomas. Treatment of Wilms’ tumor: Results of the third National Wilms’ Tumor Study. *Cancer*, 64(2):349–360, 1989. URL [https://doi.org/10.1002/1097-0142\(19890715\)64:2<349::AID-CNCR2820640202>3.0.CO;2-Q](https://doi.org/10.1002/1097-0142(19890715)64:2<349::AID-CNCR2820640202>3.0.CO;2-Q). [p4]
- V. P. Godambe. An optimum property of regular maximum likelihood estimation. *The Annals of Mathematical Statistics*, 31(4):1208–1211, 1960. URL <https://doi.org/10.1214/aoms/1177705693>. [p3]
- V. P. Godambe. Conditional likelihood and unconditional optimum estimating equations. *Biometrika*, 63(2):277–284, 1976. URL <https://doi.org/10.1093/biomet/63.2.277>. [p3]
- R. J. Gray. Weighted analyses for cohort sampling designs. *Lifetime Data Analysis*, 15(1):24–40, 2009. URL <https://doi.org/10.1007/s10985-008-9095-z>. [p4]
- D. M. Green, N. E. Breslow, J. B. Beckwith, J. Z. Finklestein, P. E. Grundym, P. R. Thomas, T. Kim, S. J. Shochat, G. M. Haase, M. L. Ritchey, P. P. Kelalis, and G. J. D’Angio. Comparison between single-dose and divided-dose administration of dactinomycin and doxorubicin for patients with Wilms’ tumor: a report from the National Wilms’ Tumor Study Group. *Journal of Clinical Oncology*, 16(1):237–256, 1998. URL <https://doi.org/10.1200/JCO.1998.16.1.237>. [p4]
- R. R. Huxley, F. L. Lopez, R. F. MacLehose, J. H. Eckfeldt, D. Couper, C. Leiendecker-Foster, R. C. Hoogeveen, L. Y. Chen, E. Z. Soliman, S. K. Agarwal, and A. Alonso. Novel association between plasma matrix metalloproteinase-9 and risk of incident atrial fibrillation in a case-cohort study: The Atherosclerosis Risk in Communities study. *PLOS One*, 8(3):1–8, 2013. URL <https://doi.org/10.1371/journal.pone.0059052>. [p5]
- InterAct Consortium. Design and cohort description of the interact project: An examination of the interaction of genetic and lifestyle factors on the incidence of type 2 diabetes in the EPIC study. *Diabetologia*, 54(9):2272–2282, 2011. URL <https://doi.org/10.1007/s00125-011-2182-9>. [p2]
- J. Jiao. *Comparison of Variance Estimators in Case-Cohort Studies*. PhD thesis, University of Southern California, 2001. URL <http://digitallibrary.usc.edu/cdm/ref/collection/p15799coll16/id/242761>. [p3, 7]
- E. Jones. *Cchs: Cox Model for Case-Cohort Data with Stratified Subcohort-Selection*, 2017. URL <https://cran.r-project.org/package=cchs>. R package version 0.4.0. [p1]
- E. Jones, M. J. Sweeting, S. J. Sharp, and S. G. Thompson. A method making fewer assumptions gave the most reliable estimates of exposure-outcome associations in stratified case-cohort studies. *Journal of Clinical Epidemiology*, 68(12):1397–1405, 2015. URL <https://doi.org/10.1016/j.jclinepi.2015.04.007>. [p2, 5]
- S. P. Juraschek, G. P. S. Shantha, A. Y. Chu, E. R. Miller III, E. Guallar, R. C. Hoogeveen, C. M. Ballantyne, F. L. Brancati, M. I. Schmidt, J. S. Pankow, and J. H. Young. Lactate and risk of incident diabetes in a case-cohort of the Atherosclerosis Risk in Communities (ARIC) study. *PLOS One*, 8(1):1–7, 2013. URL <https://doi.org/10.1371/journal.pone.0055113>. [p5]

- H. A. Katki and S. D. Mark. Survival analysis for cohorts with missing covariate information. *R News*, 8(1):14–19, 2008. URL [https://www.r-project.org/doc/Rnews/Rnews\\_2008-1.pdf#section\\*.31](https://www.r-project.org/doc/Rnews/Rnews_2008-1.pdf#section*.31). [p2]
- J. P. Kim, W. Lu, T. Sit, and Z. Ying. A unified approach to semiparametric transformation models under general biased sampling schemes. *Journal of the American Statistical Association*, 108(501): 217–227, 2013. URL <https://doi.org/10.1080/01621459.2012.746073>. [p4]
- M. Kulich and D. Lin. Improving the efficiency of relative-risk estimation in case-cohort studies. *Journal of the American Statistical Association*, 99(467):832–844, 2004. URL <https://doi.org/10.1198/016214504000000584>. [p3, 4]
- M. M. Lamb, M. D. Simpson, J. Seifert, F. W. Scott, M. Rewers, and J. M. Norris. The association between IgG4 antibodies to dietary factors, islet autoimmunity and type 1 diabetes: The diabetes autoimmunity study in the young. *PLOS One*, 8(2):1–7, 2013. URL <https://doi.org/10.1371/journal.pone.0057936>. [p5]
- B. Langholz and J. Jiao. Computational methods for case-cohort studies. *Computational Statistics & Data Analysis*, 51(8):3737–3748, 2007a. URL <https://doi.org/10.1016/j.csda.2006.12.028>. [p1, 2, 3, 4, 6, 7, 9]
- B. Langholz and J. Jiao. Case-cohort computation. <http://web.archive.org/web/20100720003853/http://hydra.usc.edu/timefactors/Examples/Case-cohort%20computational%20methods/Case-cohort%20computation.html>, 2007b. [p1]
- B. Lindsay. Conditional score functions: Some optimality results. *Biometrika*, 69(3):503–512, 1982. URL <https://doi.org/10.1093/biomet/69.3.503>. [p3]
- D. Liu, T. Cai, and Y. Zheng. Evaluating the predictive value of biomarkers with stratified case-cohort design. *Biometrics*, 68(4):1219–1227, 2012. URL <https://doi.org/10.1111/j.1541-0420.2012.01787.x>. [p4]
- T. Lumley. Analysis of complex survey samples. *Journal of Statistical Software*, 9(8):1–19, 2004. URL <https://doi.org/10.18637/jss.v009.i08>. [p2]
- T. Lumley. *Survey: Analysis of Complex Survey Samples*, 2017a. URL <https://cran.r-project.org/package=survey>. R package version 3.32-1. [p2]
- T. Lumley. *Two-Phase Designs in Epidemiology*, 2017b. URL <https://cran.r-project.org/web/packages/survey/vignettes/epi.pdf>. [p4]
- S. D. Mark and H. A. Katki. Specifying and implementing nonparametric and semiparametric survival estimators in two-stage (nested) cohort studies with missing case data. *Journal of the American Statistical Association*, 101(474):460–471, 2006. URL <https://doi.org/10.1198/016214505000000952>. [p2, 4]
- R. Payne, M. Neykov, M. K. Jensen, and T. Cai. Kernel machine testing for risk prediction with stratified case cohort studies. *Biometrics*, 72(2):372–381, 2016. URL <https://doi.org/10.1111/biom.12452>. [p4]
- R. L. Prentice. A case-cohort design for epidemiologic cohort studies and disease prevention trials. *Biometrika*, 73(1):1–11, 1986. URL <https://doi.org/10.1093/biomet/73.1.1>. [p1, 2, 3, 4]
- D. B. Rubin. *Multiple Imputation for Nonresponse in Surveys*. John Wiley & Sons, New York, 1987. URL <https://doi.org/10.1002/9780470316696>. [p5]
- S. O. Samuelsen, H. Ånestad, and A. Skrondal. Stratified case-cohort analysis of general cohort sampling designs. Technical report, Department of Mathematics, University of Oslo, 2006. URL <https://www.duo.uio.no/handle/10852/10351>. [p3]
- S. O. Samuelsen, H. Ånestad, and A. Skrondal. Stratified case-cohort analysis of general cohort sampling designs. *Scandinavian Journal of Statistics*, 34(1):103–119, 2007. URL <https://doi.org/10.1111/j.1467-9469.2006.00552.x>. [p2, 3]
- SAS Institute Inc. *SAS Software, Version 8*. Cary, NC, 1999. URL <http://www.sas.com/>. [p1]
- J. L. Schafer. Multiple imputation: a primer. *Statistical Methods in Medical Research*, 8(1):3–15, 1999. URL <https://doi.org/10.1177/096228029900800102>. [p5]

- T. A. Severini. Frequency properties of inferences based on an integrated likelihood function. *Statistica Sinica*, 21(1):433–447, 2011. URL <https://www.jstor.org/stable/24309279>. [p3]
- S. J. Sharp, M. Poulaliou, S. G. Thompson, I. R. White, and A. M. Wood. A review of published analyses of case-cohort studies and recommendations for future reporting. *PLOS One*, 9(6):e101176, 2014. URL <https://doi.org/10.1371/journal.pone.0101176>. [p5, 9]
- T. M. Therneau. *Survival: A Package for Survival Analysis in S*, 2017. URL <https://cran.r-project.org/package=survival>. R package version 2.41-3. [p2]
- T. M. Therneau and P. M. Grambsch. *Modeling Survival Data: Extending the Cox Model*. Springer-Verlag, New York, 2000. [p2]
- T. M. Therneau and H. Li. Computing the Cox model for case-cohort designs. *Lifetime Data Analysis*, 5(2):99–112, 1999. URL <https://doi.org/10.1023/A:1009691327335>. [p2, 7]
- Tibco Software Inc. *S-Plus Software, Version 8*. Palo Alto, CA, 2007. URL <http://www.tibco.com/>. [p1]
- Y. Yan, H. Zhou, and J. Cai. Improving efficiency of parameter estimation in case-cohort studies with multivariate failure time data. *Biometrics*, 73(3):1042–1052, 2017. URL <https://doi.org/10.1111/biom.12657>. [p4]
- D. Zeng and D. Y. Lin. Maximum likelihood estimation in semiparametric regression models with censored data. *Journal of the Royal Statistical Society B*, 69(4):507–564, 2007. URL <https://doi.org/10.1111/j.1369-7412.2007.00606.x>. [p4]
- Ørnulf Borgan, B. Langholz, S. O. Samuelsen, L. Goldstein, and J. Pogoda. Exposure stratified case-cohort designs. *Lifetime Data Analysis*, 6(1):39–58, 2000. URL <https://doi.org/10.1023/A:1009661900674>. [p1, 2, 3, 4, 7, 9]

Edmund Jones  
Cardiovascular Epidemiology Unit  
Department of Public Health and Primary Care  
University of Cambridge  
United Kingdom  
[edmundjones79@gmail.com](mailto:edmundjones79@gmail.com)