The New R Help System

by Duncan Murdoch and Simon Urbanek

Abstract: Version 2.10.0 of R includes new code for processing '.Rd' help files. There are some changes to what is allowed, and some new capabilities and opportunities.

One of the reasons for the success of R is that package authors are required to write documentation in a structured format for all of the functions and datasets that they make visible in their packages. This means users can count on asking for help on a topic `"foo"` using either the function call `help("foo")`, or one of the shortcuts: entering `?foo` in the console, or finding help through the menu system in one of the graphical user interfaces. The quality of the help is improved by the structured format: the quality assurance tools such as `R CMD check` can report inconsistencies between the documentation and the code, undocumented objects, and other errors, and it is possible to build indices automatically and do other computations on the text.

The original help system was motivated by the help system for S (Becker et al., 1988), and the look of the input files was loosely based on LATEX (Lamport, 1986). Perl (Wall et al., 2000) scripts transformed the input files into formatted text to display in the R console, LATEX input files to be processed into Postscript or PDF documents, and HTML files to be viewed in a web browser.

These Perl scripts were hard to maintain, and inconsistencies crept into the system: different output formats could inadvertently contain different content. Moreover, since the scripts could only look for fairly simple patterns, the quality control software had trouble detecting many errors which would slip through and result in rendering errors in the help pages, unknown to the author.

At the useR! 2008 meeting in Dortmund, one of us (Murdoch) was convinced to write a full-fledged parser for '.Rd' files. This made it into the 2.9.0 release in April, 2009, but only as part of the quality control system: and many package authors started receiving warning and error messages. Over the summer since then several members of the R Core team (including especially Brian Ripley and Kurt Hornik, as well as the authors of this article) have refined that parser, and written renderers to replace the Perl scripts, so that now all help processing is done in R code. We have also added an HTTP server to R to construct and deliver the help pages to a web browser on demand, rather than relying on static copies of the pages.

This article describes the components of the new help system.

The Parser

In the new help system, the transformation from an '.Rd' file to a '.tex', '.html' or '.txt' file is a two-step process. The `parse_Rd()` function in the **tools** package converts the '.Rd' file to an R object with class `"Rd"` representing its structure, and the renderers convert that to the target form.

Rd Syntax

The syntax of '.Rd' files handled by the parser is in some ways more complicated than the syntax of R itself. It contains LATEX-like markup, R-like sections of code, and occasionally code from other languages as well. Figure 1 shows a near-minimal example.

```
% Comments in .Rd files start with percent signs
\name{foo}
\alias{footopic}
\title{Title to Display at the Top of the Page}
\description{
  A short description of what is being documented.
}
\usage{
  foo(arg = "\n")
}
\arguments{
  \item{arg}{the first argument.}
}
\seealso{
  \code{\link{bar}}.
}
\examples{
  ## call foo then \link{bar} in a loop
  for (i in 1:10) {
    foo(1)
    bar(2)
  }
}
\keyword{example}
```

Figure 1: The contents of a simplified '.Rd' file.

Like LATEX, a comment is introduced by a percent sign (%), and markup is prefixed with a backslash (\). Braces ({}) are used to delimit the arguments to markup macros. However, notice that in the `\examples{}` section, we follow R syntax: braces are used to mark the body of the `for` loop.

In fact, there are three different modes that the parser uses when parsing the '.Rd' file. It starts out in a LATEX-like mode, where backslashes and braces are used to indicate macros and their arguments. Some macros (e.g. `\usage{}` or `\examples{}`) switch into R-like mode: macros are still recognized, but braces are used in the code, and don't necessarily indicate the arguments to macros. The third mode is mostly verbatim: it doesn't recognize any macros at

all. It is used in a few macros like \alias{}, where we might want to set an alias involving backslashes, for instance. There are further complications in that strings in quotes in R-like mode (e.g. "\n") are handled slightly differently again, so the newline is not treated as a macro, and R comments in R-like mode have their own special rules.

A good question to ask is why the syntax is so complicated. Largely this is a legacy of the earlier versions of the help system. Because there was no parsing, just a single-step transformation to the output format, special characters were handled on a case-by-case basis, and not always consistently. When writing the new parser, we wanted to minimize the number of previously correct '.Rd' files which triggered errors or were rendered incorrectly. The three-mode syntax described above is the result.

The syntax rules are fully described in Murdoch (2009), but most users should not need to know all the details. In almost all cases, the syntax is designed so that typing what you mean will give the result you want. A quick summary is as follows:

- The characters \, %, { and } have special meaning almost everywhere in an '.Rd' file.

- The backslash \ is used both to introduce macros and to escape the special meaning of the other characters.

- Unless escaped, the percent % character starts a comment. The comment is included in the parsed object, but the renderers will not display it to the user.

- Unless escaped, the braces { and } delimit the arguments to macros. In R-like or verbatim text, they need not be escaped if they balance.

- End of line (newline) characters mark the end of pieces of text, even when following a % comment.

- Other whitespace (spaces, tabs, etc.) is included as part of the text, though the renderers may remove or change it.

- In R-like text, quoted strings follow R rules: the delimiters must balance, and braces within need not. Only a few macros are recognized in R strings: \var and those related to links. Other uses of a backslash, e.g. "\n" are taken to be part of the string.

- In R-like text, R comments using # are taken to be part of the text.

- The directives #ifdef ... #endif and #ifndef ... #endif are treated as markup, so other macros must be completely nested within them or completely contain them.

There are more than 60 macros recognized by the parser, from \acronym to \verb. Each of them takes from 0 to 3 arguments in braces, and some of them take optional arguments in brackets ([]). A complete list is given in Murdoch (2009), and their interpretation is described in the *Writing R Extensions* manual (R Development Core Team, 2009).

Rd Objects

The output of the parse_Rd() function is a list with class "Rd". Currently this is mostly intended for internal use, and the only methods defined for that class are as.character.Rd() and print.Rd(). The elements of the list are the top level components of the '.Rd' file. Each element has an "RdTag" attribute labelling it as one of the three kinds of text, or a comment, or as a macro. There is a (currently internal) function in the **tools** package to display all of the tags. For example, using the help file from Figure 1, we get the following results.

```
> library(tools)
> parsed <- parse_Rd("foo.Rd")
> tools:::RdTags(parsed)

 [1] "COMMENT"       "TEXT"
 [3] "\\name"        "TEXT"
 [5] "\\alias"       "TEXT"
 [7] "\\title"       "TEXT"
 [9] "\\description" "TEXT"
[11] "\\usage"       "TEXT"
[13] "\\arguments"   "TEXT"
[15] "\\seealso"     "TEXT"
[17] "\\examples"    "TEXT"
[19] "\\keyword"     "TEXT"
```

The TEXT components between each section are simply the newlines that separate them, and can be ignored. The components labelled with macro names are themselves lists, with the same structure. For example, component 15 is the \seealso section, and it has this structure

```
> tools:::RdTags(parsed[[15]])

[1] "TEXT"   "TEXT"   "\\code" "TEXT"
```

and the \code macro within it is a list, etc. Most users will have no need to look at "Rd" objects at this level, but this is crucial for the internal quality assurance code, and for the renderers described in the next section. As R 2.10.0 is the first release where these objects are being fully used, there have been a number of changes since they were introduced in R 2.9.0, and there may still be further changes in future releases: so users of the internal structure should pay close attention to changes taking place on the R development trunk.

Incompatibilities

One of the design goals in writing the new parser was to accept valid '.Rd' files from previous versions. This was mostly achieved, but there are some incompatibilities arising from the new syntax rules given above. A full description is included in Murdoch (2009); here we point out some highlights.

In previous versions, the \code{} macro was used for both R code and code in other languages. However, the R code needs R-like parsing rules, and other languages often need verbatim parsing. Since R code is more commonly used, it was decided to restrict \code{} to handle R code and introduce \verb{} for the rest. At present, this mainly affects text containing quote marks, which must balance in a \code{} block as in R code. At some point in the future legal R syntax might be more strictly enforced.

The handling of #ifdef/#ifndef and the rules for escaping characters are now more regular.

Error Handling

One final feature of parse_Rd() is worth mentioning. We make a great effort to report errors in '.Rd' files at the point they occur, and because we now fully parse the file, we have detected a large number of errors that previously went unnoticed. In most cases these errors resulted in help pages that were not being displayed as the author intended, but we don't want to suddenly make hundreds of packages on CRAN unusable. The compromise we reached is as follows: when the parser detects an error in a '.Rd' file, it reports an error or warning, and attempts to recover, possibly skipping some text. The package installation code will report these messages, but will not abort an installation because of them. Authors should not ignore these messages: in most cases, they indicate that something will not be displayed as intended.

The Renderers

As of version 2.10.0, R will present help in three different formats: text, HTML, and PDF manuals. Previous versions also included compiled HTML on the Windows platform; that format is no longer supported, due partly to security concerns (Microsoft Support, 2007), partly to the fact that it will not support dynamic help, and partly to reduce the support load on the R Core team. The displayed type is now controlled by a single option rather than a collection of options: for example, to see HTML help, use options(help_type="html") rather than options(htmlhelp=TRUE).

The three formats are produced by the functions Rd2txt(), Rd2HTML() and Rd2latex() in the **tools** package. The main purpose of these functions is to convert the parsed "Rd" objects to the output format; they also accept '.Rd' files as input, calling parse_Rd() if necessary.

There are two other related functions in **tools**. The Rd2ex() function extracts example code from a help page, and checkRd() performs checks on the contents of an "Rd" object. The latter checks for the presence of all necessary sections (at most once in some cases, e.g. for the \title{}). It is used to report errors and warnings during R CMD check processing of a package.

One big change from earlier versions of R is that these functions do their rendering on request. In earlier versions, all help processing was done when the package was installed, and the text files, HTML files, and PDF manuals were installed with the package. Now the "Rd" objects are produced at install time, but the human-readable displays are produced at the time the user asks to see them. As described below, this means that help pages can include information computed just before the page is displayed. Not much use is made of this feature in 2.10.0, but the capability is there, and we expect to make more use of it in later releases of base R, and to see it being used in user-contributed packages even sooner.

Another advantage of "just-in-time" rendering is that cross-package links between help pages are handled better. In previous versions of R, the syntax

```
\link[stats]{weighted.mean}
```

meant that the HTML renderer should link to the *file* named 'weighted.mean.html' in the **stats** package, whereas

```
\link{weighted.mean}
```

meant to link to whatever file corresponded to the *alias* weighted.mean in the current package. This difference was necessary because the installer needed to build links at install time, but it could not necessarily look up topics outside of the current package (the target package might not be installed yet). Unfortunately, the inconsistency was a frequent source of errors, even in the base packages. Now that calculation can be delayed until rendering time, R 2.10.0 supports both kinds of linking. For back-compatibility, it gives priority to linking by filename, but if that fails, it will try to link by topic.

The HTTP server

In order to produce the help display in a web browser when the user requests, it was necessary to add an HTTP server to R.

At a high level the goal of the server is to accept connections from browsers, convert each HTTP request on a TCP/IP port into a call to an R function and deliver the result back to the browser. Since the main use is to provide a dynamic help system based

on the current state of the R session (such as packages loaded, methods defined, ...), it is important that the function is evaluated in the current R session. This makes the implementation more tricky as we could not use regular socket connections and write the server entirely in R.

Instead, a general asynchronous server infrastructure had to be created which handles multiple simultaneous connections and is yet able to synchronize them into synchronous calls of the R evaluator. The processing of each connection is dispatched to an asynchronous worker which keeps track of the state of the connection, collects the entire HTTP request and issues a message to R to process the request when complete. The request is processed by calling the `httpd` function with two arguments: path from the URL and query parameters. The query parameters are parsed from the query string into a named string vector and decoded, so for example `text=foo%3f&n=10` is passed as `c(text="foo?", n="10")`.

The `httpd()` function is expected to produce output in the form of a list that is meaningful to the browser. The list's elements are interpreted in sequence as follows: *payload*, *content-type*, *headers* and *status code*. Only the payload is mandatory so the returned list can have one to four elements. The content type specifies the MIME-type of the payload (default is "text/html"), headers specify optional HTTP response headers as a named character vector (without CR/LF) and the error code is an integer specifying the HTTP status code. The payload must be either a string vector of length one or a raw vector. If the payload element is a string vector named "file" then the string is interpreted as an absolute path to a file to be sent as the payload (useful for static content). Otherwise the payload is sent to the browser in verbatim. The `httpd` function is evaluated in a `tryCatch` block such that errors are returned as a string from the function, resulting in a 500 status code (internal server error). An example of a simple `httpd` function is listed below:

```
httpd <- function(path, query, ...) {
  if (is.null(query)) query <- character(0)
  pkg <- query['package']
  if (is.na(pkg)) pkg <- 'base'
  fn <- system.file(path, package = pkg)
  if (file.exists(fn)) return(list(file = fn))
  list(paste("Cannot find", path),
       "text/html", NULL, 404L)
}
```

The signature of the `httpd` function should include ... for future extensions. The above function checks for the `package` query parameter (defaulting to "base") then attempts to find a file given by `path` in that package. If successful the content of the file is served (as default text/html), otherwise a simple error page is created with a 404 HTTP status code "not found".

Although the HTTP server is quite general and could be used for many purposes, R 2.10.0 has the limitation of one server per R instance which makes it currently dedicated to the dynamic help. The user may set `options("help.ports")` to control which IP port is used by the server. If not set, the port is assigned randomly to avoid collisions between different R instances.

New Features in '.Rd' Files

R Expressions in Help

As mentioned previously, help output is now being produced just before display, and it is possible for the author to customize the display at that time. This is supported by the new macro \Sexpr{}, which is modelled after the macro of the same name in current versions of Sweave (Leisch, 2002).

Unlike Sweave, in '.Rd' files the \Sexpr{} macro takes optional arguments to control how it is displayed, and when it is calculated. For example, \Sexpr[results=verbatim,echo=TRUE]{x<-10;x^2} will result in a display similar to

```
> x<-10;x^2

[1] 100
```

when the help page is displayed.

Some of the options are similar to Sweave, but not always with the same defaults. For example the options (with defaults) `eval=TRUE`, `echo=FALSE`, `keep.source=TRUE`, and `strip.white=TRUE` do more or less the same things as in Sweave. That is, they control whether the code is evaluated, echoed, reformatted, and stripped of white space before display, respectively.

There are also options that are different from Sweave. The `results` option has the following possible values:

results=text (the default) The result should be inserted into the text at the current point.

results=verbatim Print the result as if it was executed at the console.

results=hide No result should be shown.

results=rd The result should be parsed as if it was '.Rd' file code.

The `stage` option controls when the code is run. It has the following values:

stage=build The code should be run when building the package. This option is not currently implemented, but is allowed; the code will not be executed in R 2.10.0.

stage=install (the default) The code is run when installing the package from source, or when building a binary version of the package.

stage=render The code is run just before the page is rendered.

The `\Sexpr{}` macro enables a lot of improvements to the help system. It will allow help pages to include examples with results inline, so that they can be mixed with descriptions, making explanations clearer. The just-in-time rendering will also allow help pages to produce information customized to a particular session: for example, it would be helpful to link from a page about a class to methods which mention it in their signatures. It will also be possible to prototype new types of summary pages or indices for the whole help system, without requiring changes to base R.

Conditional markup

The help files have supported `#ifdef`/`#ifndef` conditionals for years, to allow selective inclusion of text depending on the computing platform where R is run. With version 2.10.0 more general conditional selection has been added: the `\if{}{}` and `\ifelse{}{}{}` macros. The first argument to these describes the condition, the second is code to be included at render time if that condition holds, and the third if it does not.

The most common use of the conditionals is foreseen to be to select displays depending on the output format, to extend the idea of the `\eqn` and `\deqn` macros from previous versions. To support this use, the condition is expressed as a comma-separated list of formats, chosen from `example`, `html`, `latex`, `text`, `TRUE` and `FALSE`. If the current output format or `TRUE` is in the list, the condition is satisfied. The logicals would normally be the result of evaluating an `\Sexpr{}` expression, so general run-time conditions are possible.

Verbatim output

Two new macros have been added to support output of literal text. The `\verb{}` macro parses its argument as verbatim text, and the renderers output it as parsed, inserting whatever escapes and conversions are necessary so that it renders the same in all formats. The `\out{}` macro works at a lower level. It also parses its argument as verbatim text, but the renderers output it in raw form, without any processing. It would almost certainly be used in combination with `\if` or `\ifelse`. For example, to output a Greek letter alpha, the markup

```
\ifelse{html}{\out{&alpha;}}{
\eqn{\alpha}{alpha}}
```

could be used: this will output `α` when producing HTML and `\alpha` when producing LaTeX, both of which will render as α; it will output `alpha` when producing text.

The Future

One of the development guidelines for R is that we don't introduce many new features in patch releases, so R 2.10.1 is not likely to differ much from what is described above. However, we would expect small changes to the rendering of help pages as the renderers become more polished.

For version 2.11.0 or later, we would expect some of the following to be implemented. The order of presentation is not necessarily the order in which they will appear, and some may never appear.

- The `\Sexpr{}` macro will likely develop in several ways:

  - We will likely add in the processing of `stage=build` macros when the `R CMD build` code is rewritten in R.

  - We will likely change the `prompt()`, `package.skeleton()` and related functions to make use of `\Sexpr{}` macros, e.g. to display the version of a package, or to list the classes supporting a method, etc.

  - We will make information about the run-time environment available to the `\Sexpr{}` code, so that it can tailor its behaviour to the context in which it is being run.

  - We would expect users to develop contributed packages to provide convenient functions to use in `\Sexpr{}` macros. Some of these may need changes to base R.

  - We would like to allow figures to be inserted into help displays.

- The quality control checks will continue to evolve, as we notice common errors and add code to `checkRd()` to detect them. There is already support for spell checking of '.Rd' files in the `aspell()` function.

- The '.Rd' format is very idiosyncratic, and not many tools exist outside of R for working with it. There may be an advantage to switching to a different input format (e.g. one based on XML) in order to make use of more standard tools. Such a change would require conversion of the thousands of existing help pages already in '.Rd' format; the new parser may enable automatic translation if anyone wants to explore this possibility.

- The HTTP server used in the help system is quite general, but is currently limited to serving R help pages. It may be extended to allow more general use in the future.

Changing to the new help system has already helped to diagnose hundreds of errors in help pages that slipped by in the previous version, and diagnostic messages are more informative than they were. As the `\Sexpr{}` macro is used in core R and by package writers, we will see better help than ever before, and the HTTP server will open up many other possibilities for R developers.

Bibliography

R. A. Becker, J. M. Chambers, and A. R. Wilks. *The New S Language*. Wadsworth, Pacific Grove, California, 1988.

L. Lamport. *LaTeX: A Document Preparation System*. Addison-Wesley, Reading, Massachusetts, 1986.

F. Leisch. Sweave: Dynamic generation of statistical reports using literate data analysis. In W. Härdle and B. Rönz, editors, *Compstat 2002 — Proceedings in Computational Statistics*, pages 575–580. Physica Verlag, Heidelberg, 2002. URL http://www.stat.uni-muenchen.de/~leisch/Sweave. ISBN 3-7908-1517-9.

Microsoft Support. Ms05-026: A vulnerability in html help could allow remote code execution, 2007. Retrieved from http://support.microsoft.com/kb/896358 on November 5, 2009.

D. Murdoch. Parsing Rd files. 2009. URL http://developer.r-project.org/parseRd.pdf.

R Development Core Team. *Writing R Extensions*, 2009. Manual included with R version 2.10.0.

L. Wall, T. Christiansen, and J. Orwant. *Programming Perl*. O'Reilly Media, third edition, 2000.

Duncan Murdoch
Dept. of Statistical and Actuarial Sciences
University of Western Ontario
London, Ontario, Canada
murdoch@stats.uwo.ca

Simon Urbanek
AT&T Labs – Statistics Research
Florham Park, NJ, USA
urbanek@research.att.com